

## Problems

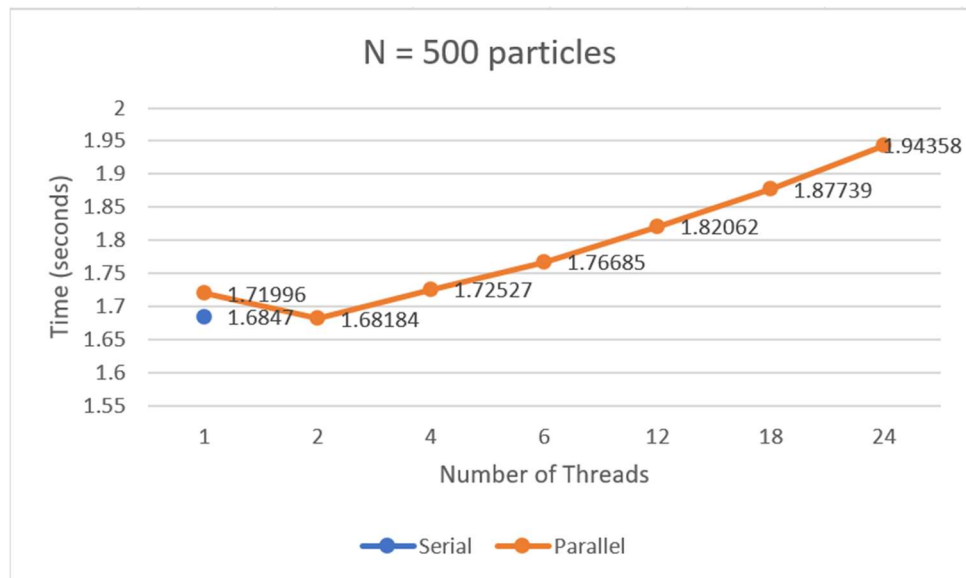
- 1) Code segment A is a simple parallelized program, however this does not have any included mutual exclusion, so issues could arise when running this on multiple threads accessing shared memory. B and C both include mutual exclusion. In this scenario, both B and C should perform the same, however their functionality is different in that code block B is allowing for threads to lock the array x itself while code block C will have a separate lock that will be used on the entire code block following the critical statement, which in this case is just the single line using array x.
- 2)
  - a. On average, each thread should take  $2i$  ms to complete. When a thread has to wait for the other, it should only wait for  $i$  ms.
  - b. The execution time should not change.
  - c. Adding dynamic has the possibility to decrease out execution time if the threads become out of sync, due to the dynamic assignment nature. But in worst-case, the execution time may substantially increase due to wait times and added overhead.
  - d. Using firstprivate or lastprivate will help in this scenario.

## Programming Assignment

- 3) Programming Assignment

\*\*\*My files are located under '/ihome/ece2192-2022f/cum6/ece1570/hw2'

```
n = 500, simulation time = 1.6847 seconds
n = 500, threads = 1, simulation time = 1.71996 seconds
n = 500, threads = 2, simulation time = 1.68184 seconds
n = 500, threads = 4, simulation time = 1.72527 seconds
n = 500, threads = 6, simulation time = 1.76685 seconds
n = 500, threads = 12, simulation time = 1.82062 seconds
n = 500, threads = 18, simulation time = 1.87739 seconds
n = 500, threads = 24, simulation time = 1.94358 seconds
n = 1000, threads = 2, simulation time = 4.15956 seconds
n = 2000, threads = 4, simulation time = 7.3319 seconds
n = 3000, threads = 6, simulation time = 10.7801 seconds
n = 6000, threads = 12, simulation time = 26.3469 seconds
n = 9000, threads = 18, simulation time = 39.3043 seconds
n = 12000, threads = 24, simulation time = 53.6345 seconds
```



My methodology was simply to break down the entire particle dispersal area into a series of bins. Then once we have placed all these particles into bins, we can calculate the forces from the neighboring bins. When parallelizing this, I made sure to include barriers between each phase to ensure that threads are at the same point.