

Outline

- Processors and registers
- Memory hierarchies
 - Temporal and spatial locality
 - Cache timing
 - Use of microbenchmarks to characterize performance
- Case study: Matrix multiplication
- Optimizations in Practice

Idealized Uniprocessor Model

- Processor names variables:
 - Integers, floats, doubles, pointers, arrays, structures, ... are really bytes, words, etc. in address space
- Processor performs operations on those variables:
 - Arithmetic operations, etc. only on values in registers
 - Load/Store variables between memory and registers
- Processor controls the order, as specified by program
 - Branches (if), loops, function calls, etc.
 - Compiler translates into “obvious” lower level instructions
 - Hardware executes instructions in order specified by compiler
 - Read returns the most recently written data
- Idealized Cost
 - Each operation (+,*, &, etc.) has roughly the same cost (on “free” registers)
 - Load/store is 100x the cost of +,*,&, etc.

Compilers Manage Memory and Registers

- Compiler performs “register allocation” to decide when to load/store vs reuse

a R1

b R4

f R1

Register allocation in first Fortran compiler in 1950s, graph coloring in 1980.

$a = c + d$

$e = a + b$

$f = e - 1$

Assume a and e not used again (dead), b,c,d, and f are live

Load c into R2

Load d into R3

Load b into R4

$R1 = R2 + R3$

$R1 = R1 + R4$

$R1 = R1 - 1$


Optimizing Compilers

- Besides register allocation, the compiler performs optimizations:
 - Unrolls loops (because control isn't free)
 - Fuses loops (merge two together)
 - Interchanges iteration variables (reorder, can lead to vectorization)
 - Eliminates dead code (the branch never taken)
 - Reorders instructions to improve register reuse and more
 - Strength reduction (e.g., shift left rather than multiply by 2)
- Why should you care?
 - Because sometimes it does the best thing possible
 - But other times it does not...

More Realistic Uniprocessor Model


- Memory accesses have *two* costs:
 - Latency (e.g. 10s to 100 ns)
 - Bandwidth: Average rate in bytes/second (or, inverse, seconds/byte) : (2-40 GB/s → 0.025 to 0.5 ns/byte)
- Bandwidth can't keep up with processor demands
- Latency is worse: DRAM isn't getting faster

Memory Bandwidth is falling behind




Memory Latency is worse

Memory Latency is much worse: $(\text{GFLOP/s}) / (\text{Memory Latency})$




Interconnect Bandwidth is also falling behind

Interconnect Bandwidth is Falling Behind at a comparable rate



Interconnect Latency follows...

Interconnect latency follows a similar trend...



Why is memory bandwidth slowly increasing?

- Slow rate of pin speed increase
 - RAMBUS™ is an exception
 - Emphasis on capacity, not speed
 - Multiplexed DIMMs
- 6-T SRAM cell time unchanged
- Pins cost money (and manufacturers are cheap)

Why is memory latency slowly increasing?

- More levels in cache hierarchy
- Many different clock domains (Intel example)
 - Local Memory: “Core” → Ring → DDR → Ring → “Core”
 - Snoopy: (n.b. QPI = QuickPath Interconnect)
“Core” → Ring → QPI → Ring → QPI → Ring → “Core”
- More cores to keep coherent (and coherency takes time)

Memory Hierarchy

- Most programs have a high degree of locality
 - spatial locality: accessing things nearby previous accesses
 - temporal locality: reusing an item that was previously accessed
- Caches may make latency (worst case) even worse
 - Robot loading tapes at supercomputer center
- Latest change: “nonvolatile memory” (NVM), like Flash
 - Flash: Storage without power; Reading and writing need power. But writing can take 2x – 10x longer than reading.

Why Have Multiple Levels of Cache?

- On-chip vs. off-chip
 - On-chip caches are faster, but smaller
- A large cache has delays
 - Hardware to check longer addresses in cache takes more time
 - Associativity, which gives a more general set of data in cache, also takes more time
- Some examples:
 - Cray T3E eliminated one cache to speed up misses
 - Intel Haswell uses a Level 4 cache as “victim cache”
- There are other levels of the memory hierarchy
 - Register, pages (TLB, virtual memory), ...
 - And it isn't always a hierarchy

Approaches to Handling Memory Latency

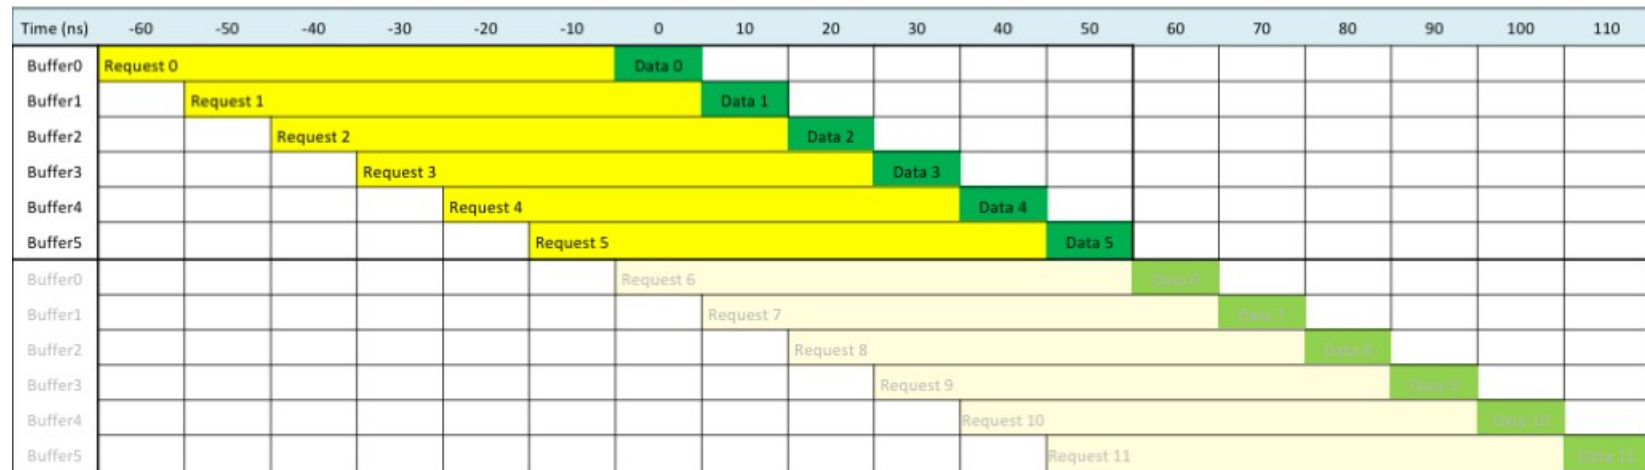
- Reuse values in fast memory (bandwidth filtering)
 - need temporal locality in program
- Move larger chunks (achieve higher bandwidth)
 - need spatial locality in program
- Issue multiple reads/writes in single instruction (higher bandwidth)
 - vector operations require access set of locations (typically neighboring)
- Issue multiple reads/writes in parallel (hide latency)
 - *prefetching* issues read hint
 - *delayed writes* (write buffering) stages writes for later operation
 - both require that nothing dependent is happening (parallelism)

Latency, Bandwidth and Concurrency

- *Little's Law* from queuing theory says:
 $\text{concurrency} = \text{latency} * \text{bandwidth}$
- For example:
 latency = 10 sec
 bandwidth = 2 Bytes/sec
- Requires 20 bytes in flight to hit bandwidth speeds
- That means finding 20 independent things to issue

Real example

Little's Law: illustration for 2005-era Opteron processor
60 ns latency, 6.4 GB/s (=10ns per 64B cache line)



- $60 \text{ ns} * 6.4 \text{ GB/s} = 384 \text{ Bytes} = 6 \text{ cache lines}$
- To keep the pipeline full, there must always be 6 cache lines “in flight”
- Each request must be launched at least 60 ns before the data is needed

Memory Benchmark (CacheBench)

for vl = all vector lengths

memory[vl] //alloc+init

timer start

for iteration count

for i=0 to vl

register r += memory[i]

timer stop

Why Matrix Multiplication?

- An important kernel in many problems
 - Dense linear algebra is extremely common
 - Closely related to other algorithms, e.g., transitive closure on a graph
 - And dominates training time in deep learning (CNNs)
- Optimization ideas can be used in other problems
- The best case for optimization payoffs
- The most-studied algorithm in high performance computing

A Simple Model of Memory

- Assume just 2 levels in the hierarchy: fast and slow
- All data initially in slow memory
 - m = number of memory elements (words) moved between fast and slow memory
 - t_m = time per slow memory operation (inverse bandwidth in best case)
 - f = number of arithmetic operations
 - t_f = time per arithmetic operation $\ll t_m$
 - Computational Intensity (CI) = f / m average number of flops per slow memory access
- Minimum possible time: $f * t_f$ when all data in fast memory
- Actual time: $f * t_f + m * t_m = f * t_f * (1 + t_m/t_f * 1/CI)$
- Larger CI means time closer to minimum $f * t_f$

Matrix-Vector Multiplication

```
// implements  $y = y + A \cdot x$   
for i = 1:n  
    for j = 1:n  
         $y(i) = y(i) + A(i,j) \cdot x(j)$ 
```

Matrix-Vector Multiplication (with memory)

```
// implements  $y = y + A \cdot x$   
// read  $x(1:n)$  into fast memory  
// read  $y(1:n)$  into fast memory  
for i = 1:n  
    // read row i of A into fast memory  
    for j = 1:n  
         $y(i) = y(i) + A(i,j) \cdot x(j)$   
    end  
// write  $y(1:n)$  back to slow memory
```

Matrix-Vector Multiplication (complexity)

- Assume for simplicity vectors x and y , and one row of A , all fit in fast memory
- m = number of slow memory refs = $3n + n^2$
- f = number of arithmetic operations = $2n^2$
- $q = f / m \approx 2$ (Low Computational Intensity)
- Matrix-vector multiplication limited by slow memory speed

Simplifying Assumptions

What simplifying assumptions did we make in this analysis?

- Constant “peak” computation rate
- Fast memory was large enough to hold three vectors
- The cost of a fast memory access is 0
 - OK for registers
 - Not for cache (even L1)
- Memory latency is constant
- Could simplify further by ignoring memory operations on x & y
 - Not bad: time to read matrix (bandwidth) may dominate

Naïve Matrix Multiply

```
// implements  $C = C + A*B$ 
```

```
for i = 1 to n
```

```
    for j = 1 to n
```

```
        for k = 1 to n
```

```
             $C(i,j) = C(i,j) + A(i,k) * B(k,j)$ 
```

- Algorithm has $2*n^3 = O(n^3)$ Flops and operates on $3*n^2$ words of memory
- Computational intensity (q) potentially as large as $2*n^3 / 3*n^2 = O(n)$

Naïve Matrix Multiply

// implements $C = C + A*B$

for $i = 1$ to n

 for $j = 1$ to n

 for $k = 1$ to n

$C(i,j) = C(i,j) + A(i,k) * B(k,j)$

of slow memory ops:

$m = n^3$ to read each column of B n times

+ n^2 to read each row of A once

+ $2n^2$ to read and write each element of C once

$= n^3 + 3n^2$

So $q = f / m = 2n^3 / (n^3 + 3n^2)$ computational intensity

≈ 2 for large n , **no** improvement over matrix-vector multiply

Blocked (Tiled) Matrix Multiply

```
// implements C = C + A*B
// A,B,C to be N-by-N matrices of b-by-b subblocks
for i = 1 to n
  for j = 1 to n
    for k = 1 to n
      // Matrix multiply over blocks of size b = n / N
      C(i,j) = C(i,j) + A(i,k) * B(k,j)
```

if $n \% b \neq 0$ you need to have code for the “edges/fringes”

$2n^2$ to read and write each block of C once: $2N^2 * b^2 = 2n^2$

$N * n^2$ to read each block of A N^3 times: $N^3 * b^2 = N^3 * (n/N)^2$

$N * n^2$ to read each block of B N^3 times: $N^3 * b^2 = N^3 * (n/N)^2$

Computational Intensity, $CI = f / m = 2n^3 / ((2N + 2) * n^2) \approx n / N = b$ for large n

Tiling for Registers

If the block dimension b is very small (e.g., 2) then:

$C[i,j] = C[i,j] + A[i,k] * B[k,j]$ //do a matrix multiply on blocks

1) we can write the inner matmul without loops (unrolled)

2) The values may all fit in registers

$$C[0,0] += A[0,0] \cdot B[0,0] + A[0,1] \cdot B[1,0]$$

$$C[0,1] += A[0,0] \cdot B[0,1] + A[0,1] \cdot B[1,1]$$

$$C[1,0] += A[1,0] \cdot B[0,0] + A[1,1] \cdot B[1,0]$$

$$C[1,1] += A[1,0] \cdot B[0,1] + A[1,1] \cdot B[1,1]$$

Are we getting the same answer?

Rather than something like

$$C[i,j] += AB + AB + AB \dots AB$$

Which hardware executes as

$$C[i,j] += (((AB + AB) + AB) \dots AB)$$

The tiled version does something like

$$C[i,j] += AB + AB; \dots C[i,j] += AB + AB$$

Which has more stores to memory and is more like

$$C[i,j] += ((AB + AB) + (AB + AB))$$

Assumes + is associative:

$$(x + y) + z = x + (y + z)$$

And extra stores don't matter

Isn't strictly true for floating point, but close enough for matrix multiply!

Recursive Matrix Multiplication

Define $C = \text{RMM}(A, B, n)$

if ($n=1$) {

$C_{00} = A_{00} * B_{00};$

} else {

$C_{00} = \text{RMM}(A_{00}, B_{00}, n/2) + \text{RMM}(A_{01}, B_{10}, n/2)$

$C_{01} = \text{RMM}(A_{00}, B_{01}, n/2) + \text{RMM}(A_{01}, B_{11}, n/2)$

$C_{10} = \text{RMM}(A_{10}, B_{00}, n/2) + \text{RMM}(A_{11}, B_{10}, n/2)$

$C_{11} = \text{RMM}(A_{11}, B_{01}, n/2) + \text{RMM}(A_{11}, B_{11}, n/2)$

}

return C

Experience with Cache-Oblivious Algorithms

- Recursive Matrix Multiply is cache oblivious
- In practice, need to cut off recursion well before 1×1 blocks
- Call “micro-kernel” on small blocks
- Careful attention to micro-kernel is needed
- Pingali et al report that they never got more than $2/3$ of peak.
- Issues with Cache Oblivious (recursive) approach

Basic Linear Algebra Subroutines (BLAS)

- Industry standard interface (evolving)
- www.netlib.org/blas, www.netlib.org/blas/blast--forum
- Vendors, others supply optimized implementations
- History
 - BLAS1 (1970s): 15 different operations
 - vector operations: dot product, saxpy ($y = \alpha * x + y$), root-sum-squared, etc
 - $m = 2 * n$, $f = 2 * n$, $q = f / m = \text{computational intensity} \sim 1$ or less
 - BLAS2 (mid 1980s): 25 different operations
 - matrix-vector operations: matrix vector multiply, etc
 - $m = n^2$, $f = 2 * n^2$, $q \sim 2$, less overhead
 - somewhat faster than BLAS1
 - BLAS3 (late 1980s): 9 different operations
 - matrix-matrix operations: matrix matrix multiply, etc
 - $m \leq 3n^2$, $f = O(n^3)$, so $q = f / m$ can possibly be as large as n , so BLAS3 is potentially much faster than BLAS2
- Good algorithms use BLAS3 when possible (LAPACK & ScaLAPACK)
 - See www.netlib.org/{lapack,scalapack}
 - More later in the course

Take-Aways

- Matrix vector and matrix matrix multiplication key (linear algebra)
- Matrix vector multiplication
 - Opportunities for better parallelism and use special instructions, Fused Multiply Add, etc.
 - Limited by bandwidth ($O(2n^2)$ flops on $O(n^2)$ data)
- Matrix matrix multiplication
 - Can improve computational intensity $O(2n^3)$ flops on $O(3n^2)$ data
 - Tiling (aka blocking)
- Optimizations in practice
 - Expose parallelism to compiler (SIMD, ILP, prefetching)
 - Help compiler with register use
 - Help hardware (data layout) with cache locality (temporal and spatial)
- Optimized libraries (BLAS) exist

Optimizing in Practice

- Tiling for registers
 - loop unrolling, use of named “register” variables
- Tiling for multiple levels of cache and TLB
- Exploiting fine-grained parallelism in processor
 - superscalar; pipelining
- Complicated compiler interactions (flags)
- Hard to do by hand (but you’ll try)
- Automatic optimization an active research area

Note on Matrix Storage

- A matrix is a 2-D array of elements, but memory addresses are “1-D”
- Conventions for matrix layout
 - by column, or “column major” (Fortran default); $A(i,j)$ at $A+i+j*n$
 - by row, or “row major” (C default) $A(i,j)$ at $A+i*n+j$
 - recursive (later)
- Intense interaction with cache line size

Tips on Tuning

“We should forget bout small efficiencies, say about 97% of the time: premature optimization is the root of all evil.”– C.A.R. Hoare (quoted by Donald Knuth)

- Tradeoff: speed vs readability, debuggability, maintainability...
- Only optimize when needful
- Go for low-hanging fruit first: data layouts, libraries, compiler flags
- Concentrate on the bottleneck
- Concentrate on inner loops
- Get correctness (and a test framework) first

Tip #1: Tools & Libraries

- We have gcc. The Intel compilers are better.
- Fortran compilers often do better than C compilers (less pointer aliasing)
- Intel VTune, cachegrind, and Shark can provide useful profiling information (including information about cache misses)
- Libraries build on someone else's hard work

Tip #2: Compiler flags

- -O3: Aggressive optimization
- -march=core2: Tune for specific architecture
- -ftree-vectorize: Automatic use of SSE (supposedly)
- -funroll-loops: Loop unrolling
- -ffast-math: Unsafe floating point optimizations

Tip #3: Memory layout

- Arrange data for unit stride access
- Arrange algorithm for unit stride access
- Tile for multiple levels of cache
- Tile for registers (loop unrolling + “register” variables)

Tip #4: Use small data structures

- Smaller data types are faster
- Bit arrays vs int arrays for flags?
- Minimize indexing data — store data in blocks
- Some advantages to mixed precision calculation (float for large data structure, double for local calculation)
- Sometimes recomputing is faster than saving

Tip #5: Inline judiciously

- Function call overhead is often minor...
- ... but structure matters to optimizer!
- C++ has inline keyword to indicate inlined functions