

# Outline

- Bureaucracy: Timeline, Homework
- Review: Pthreads, OpenMP
- Synchronization
- Parallelizing loops
- Optimizations in Practice

# Synchronization

- High level synchronization included in the common core (the full OpenMP specification has *many* more):
  - Critical section
  - Barrier

# Synchronization: critical section

```
float res;
```

```
#pragma omp parallel
```

```
{    float B;  int i, id, nthrds;
```

```
    id = omp_get_thread_num();
```

```
    nthrds = omp_get_num_threads();
```

```
    for(i=id;i<niters;i+=nthrds){
```

```
        B = big_job(i);
```

```
#pragma omp critical
```

```
    res += consume (B); // Thread wait: One at a time...
```

```
    }
```

```
}
```

# Synchronization: Barrier

- Barrier: a point in a program all threads must reach before any threads are allowed to proceed.
- It is a “stand alone” pragma meaning it is not associated with user code ... it is an executable statement.

```
double Arr[8], Brr[8];
int numthrds;
omp_set_num_threads(8)
#pragma omp parallel
{   int id, nthrds;
    id = omp_get_thread_num();
    nthrds = omp_get_num_threads();
    if (id==0) numthrds = nthrds;
    Arr[id] = big_ugly_calc(id, nthrds);
#pragma omp barrier
    Brr[id] = really_big_and_ugly(id, nthrds, Arr);
}
```

# Synchronization: Using a critical section to remove impact of false sharing

```
#include <omp.h>
static long num_steps = 100000;    double step;
#define NUM_THREADS 4
void main ()
{   int nthreads; double pi=0.0;
    step = 1.0/(double) num_steps;
    omp_set_num_threads(NUM_THREADS);
#pragma omp parallel
    {
        int i, id, nthrds;  double x, sum; // scalar sum per thread
        id = omp_get_thread_num();
            nthrds = omp_get_num_threads();
            if (id == 0) nthreads = nthrds;
        for (i=id, sum=0.0; i< num_steps; i=i+nthrds) {
            x = (i+0.5)*step;
            sum += 4.0/(1.0+x*x); // no array, no sharing
        }
        #pragma omp critical
        pi += sum * step; // critical section avoids conflicts during updates
    }
}
```

# Parallelizing loops

- OpenMP easily parallelizes loops
  - Easiest when: No data dependencies (reads/write or write/write pairs) between iterations!
- Preprocessor and runtime calculate loop bounds for each thread directly from serial source
- The loop **#pragma omp for** construct splits up loop iterations among the threads in a team

```
#pragma omp parallel
```

```
{
```

```
#pragma omp for
```

```
    for (I=0;I<N;I++){ // I is made private to each thread
```

```
        big_ugly_calc(I);
```

```
    } // All threads wait here before proceeding
```

```
}
```

# Parallelizing loops (this time with feeling)

- Sequential: `for(i=0;i<N;i++) { a[i] = a[i] + b[i];}`

- Parallel region:

```
#pragma omp parallel
```

```
{  
    int id, i, Nthrds, istart, iend;  
    id = omp_get_thread_num();  
    Nthrds = omp_get_num_threads();  
    istart = id * N / Nthrds;  
    iend = (id+1) * N / Nthrds;  
    if (id == Nthrds-1) iend = N;    for(i=istart;i<iend;i++) { a[i] = a[i] + b[i];}  
}
```

- Parallel with for: (could put on the same line)

```
#pragma omp parallel
```

```
#pragma omp for
```

```
    for(i=0;i<N;i++) { a[i] = a[i] + b[i];}
```

# Loop scheduling

- `schedule` clause determines how loop iterations are divided among the thread team; no one best way
- `static([chunk])` divides iterations statically between threads (default if no hint)
- Each thread receives `[chunk]` iterations, rounding as necessary to account for all iterations
- Default `[chunk]` is `ceil( # iterations / # threads )`
- `dynamic([chunk])` allocates `[chunk]` iterations per thread, allocating an additional `[chunk]` iterations when a thread finishes
- Forms a logical work queue, consisting of all loop iterations
- Default `[chunk]` is 1
- `guided([chunk])` allocates dynamically, but `[chunk]` is exponentially reduced with each allocation



# Working with loops

- Basic approach
  - Find compute intensive loops
  - Make the loop iterations independent ... So they can safely execute in any order without loop-carried dependencies
  - Place the appropriate OpenMP directive and test

```
int i, j, A[MAX];  
    j = 5;  
    for (i=0; i< MAX; i++) {  
        j +=2;  
        A[i] = big(j);  
    }
```

```
int i, A[MAX];  
#pragma omp parallel for  
for (i=0; i< MAX; i++) {  
    int j = 5 + 2*(i+1);  
    A[i] = big(j);  
}
```

# Static vs. Dynamic

- Suppose you are doing block matrix multiplication, i.e. you split each matrix into  $B^2$  submatrices, each of size  $N/B$ -by- $N/B$  and will do  $B^3$  submatrix multiplications ( $B$  is relatively small, say  $\sim 10$ )
- However, matrices are sparse and submatrices have varying number of nonzeros (this is a real use case).

```
#pragma omp parallel for
for (int i=0; i< B; i++)
{
    for (int j=0; j< B; j++)
    {
        for (int k=0; k< B; k++)
        {
            SpGEMM(A(i,k), B(k,j), C(i,j)); // Sparse Matrix Multiply
        }
    }
}
```

# Static vs. Dynamic (again)

- Why did we use collapse? Without it, OpenMP only parallelizes the outermost loop
- We could put nested omp parallel constructs but that can create too much parallelism generation overhead
- Instead, **collapse(2)** says to the compiler to treat the whole  $B \times B$  iteration space of the first two loops (following the program) as a single contiguous iteration space to divide up

```
#pragma omp parallel for collapse(2)
for (int i=0; i< B; i++)
{
    for (int j=0; j< B; j++)
    {
        for (int k=0; k< B; k++)
        {
            SpGEMM(A(i,k), B(k,j), C(i,j));
        }
    }
}
```

# Static vs. Dynamic (again...)

- If chunks are very large and have variable work; there isn't much any schedule can do for you.
- But if chunks are rather small and have variable work, you should try dynamic

```
#pragma omp parallel schedule(dynamic) for collapse(2)
for (int i=0;i< B; i++)
{
    for (int j=0;j< B; j++)
    {
        for (int k=0;k< B; k++)
        {
            SpGEMM(A(i,k), B(k,j), C(i,j));
        }
    }
}
```

# Reduction

- We are combining values into a single accumulation variable ... there is a true dependence between loop iterations that can't be trivially removed
- This is a very common ... it is called a “reduction”.
- Support for reduction operations is included in most parallel programming environments.

```
double ave=0.0, A[MAX];   int i;
    for (i=0;i< MAX; i++) {
        ave + = A[i];
    }
    ave = ave/MAX;
```

# Reduction

- OpenMP reduction clause: **reduction** (op : list)
  - Inside a parallel construct:
  - A local copy of each list variable is made and initialized depending on the “op” (e.g. 0 for “+”)
  - Updates occur on the *local copy*
  - Local copies are reduced into a single value and combined with the original global value.
- The variables in “list” must be shared in the enclosing parallel region.

```
double ave=0.0, A[MAX];  int i;
```

```
#pragma omp parallel for reduction (+:ave)
```

```
for (i=0;i< MAX; i++) {
```

```
    ave + = A[i];
```

```
} // implicit barrier here
```

```
ave = ave/MAX;
```

# OpenMP: Reduction operands initial-values

- Many different associative operands can be used with reduction:
- Initial values are the ones that make sense mathematically.

Operator	Initial Value
+	0
*	1
-	0
min	Largest pos number
max	Largest negative number
&	~0
	0
^	0
&&	1
	0

# Pi with a loop and a reduction

```
#include <omp.h>

static long num_steps = 100000;      double step;

void main ()
{   int i;      double x, pi, sum = 0.0;
    step = 1.0/(double) num_steps;
    #pragma omp parallel
    {
        double x;
        #pragma omp for reduction(+:sum)
        for (i=0;i< num_steps; i++){
            x = (i+0.5)*step;
            sum = sum + 4.0/(1.0+x*x);
        }
    }
    pi = step * sum;
}
```



# The nowait clause

- Implicit barriers exist at the end of most (but not all) OpenMP constructs
- Barriers cause waiting (expensive). You need to understand when they are implied and how to skip them when its safe to do so.

```
double A[big], B[big], C[big];
```

```
#pragma omp parallel
```

```
{
```

```
    int id=omp_get_thread_num();
```

```
    A[id] = big_calc1(id);
```

```
#pragma omp barrier
```

```
#pragma omp for
```

```
    for(i=0;i<N;i++){C[i]=big_calc3(i,A);} // Implicit barrier here
```

```
#pragma omp for nowait
```

```
    for(i=0;i<N;i++){ B[i]=big_calc2(C, i); } // No barrier here
```

```
    A[id] = big_calc4(id);
```

```
} // Implicit barrier here due to parallel
```

# Memory model

- Shared memory programming model:
  - Most variables are shared by default
- Global variables are **shared** among threads
  - Fortran: COMMON blocks, SAVE variables, MODULE variables
  - C: Globally scoped variables, static
  - Both languages: dynamically allocated memory (ALLOCATE, malloc, new)
- But not everything is shared...
  - Stack variables in subprograms(Fortran) or functions(C) called from parallel regions are **private**
  - Automatic variables within a statement block are **private**

# Data Sharing

```
double A[10];  
int main() {  
    int index[10];  
    #pragma omp parallel  
        work(index);  
    printf("%d\n", index[0]);  
}
```

```
extern double A[10];  
void work(int *index) {  
    double temp[10];  
    static int count;  
    ...  
}
```

- A, index and count are shared by all threads.
- temp is local to each thread

# Data sharing: Attributes

- One can selectively change storage attributes for constructs using the following clauses (note: list is a comma-separated list of variables)
  - `shared(list)`
  - `private(list)`
  - `firstprivate(list)`
- These can be used on parallel and for constructs ... other than *shared* which can only be used on a parallel construct
- Force the programmer to explicitly define storage attributes
  - default (none)

# Data sharing: private

- **private(var)** creates a new local copy of var for each thread.
- The value of the private copies are *uninitialized*
- The value of the original variable is unchanged after the region

```
void wrong() {  
    int tmp = 0;  
    #pragma omp parallel for private(tmp)  
    for (int j = 0; j < 1000; ++j)  
        tmp += j; // tmp wasn't initialized here  
    printf("%d\n", tmp); // tmp is 0 here! (outside of for scope)  
}
```

# Data sharing: firstprivate

- Variables initialized from a shared variable
- C++ objects are copy-constructed

```
incr = 0;
```

```
#pragma omp parallel for firstprivate(incr)
```

```
for (i = 0; i <= MAX; i++) { // Each thread gets own incr with initial 0
```

```
    if ((i%2)==0) incr++;
```

```
    A[i] = incr;
```

```
}
```

# Data sharing: Summary

- Are A,B,C private to each thread or shared inside the parallel region?
- What are their initial values inside and values after the parallel region?

variables:  $A = 1$ ,  $B = 1$ ,  $C = 1$

`#pragma omp parallel private(B) firstprivate(C)`

Inside this parallel region ...

- “A” is shared by all threads; equals 1
- “B” and “C” are private to each thread.
  - B’s initial value is undefined
  - C’s initial value equals 1

Following the parallel region ...

- B and C revert to their original values of 1
- A is either 1 or the value it was set to inside the parallel region

# Data sharing: default

- `default(none)`: Forces you to define the storage attributes for variables that appear inside the static extent of the construct ... if you fail the compiler will complain. Good programming practice!
- The compiler would complain about `j` and `y`, which is important since you don't want `j` to be shared

```
#include <omp.h>
int main()
{
    int i, j=5;    double x=1.0, y=42.0;
    #pragma omp parallel for default(none) reduction(*:x)
    for (i=0;i<N;i++){
        for(j=0; j<3; j++)
            x+= foobar(i, j, y);
    }
    printf(" x is %f\n",(float)x);
}
```



# Tasks

- Tasks are *independent* units of work
- Tasks are composed of:
  - code to execute
  - data to compute with
- Threads are assigned to perform the work of each task.
- The thread that encounters the **task** construct may execute the task immediately.
- The threads may defer execution until later

# Tasks

- The task construct includes a structured block of code
- Inside a parallel region, a thread encountering a task construct will package up the code block and its data for execution
- Tasks can be nested: i.e. a task may itself generate tasks.

# Tasks: single

- The **single** construct denotes a block of code that is executed by only one thread (not necessarily the master thread).
- A barrier is implied at the end of the **single** block (can remove the barrier with a nowait clause).

```
#pragma omp parallel
{
    do_many_things();
#pragma omp single
    {   exchange_boundaries();   }
    do_many_other_things();
}
```

# Tasks: task directive

`#pragma omp task [clauses] structured-block`

```
#pragma omp parallel
{
    #pragma omp single
    {
        #pragma omp task
        fred();
        #pragma omp task
        daisy();
        #pragma omp task
        billy();
    } // All tasks complete before this barrier is released
}
```

# Tasks: When are they completed?

- At thread barriers (explicit or implicit)
  - applies to all tasks generated in the current parallel region up to the barrier
- At **taskwait** directive
  - i.e. wait until all tasks defined within the scope of the current task have completed.
  - `#pragma omp taskwait`
  - Note: applies only to tasks *generated in the current task, not to “descendants”* .
  - To also wait for descendants, there is the taskgroup `*region*`
  - When a thread encounters a taskwait construct, the current task is suspended until all child tasks that it generated before the taskwait region complete execution

# Tasks: task directive

`#pragma omp task [clauses] structured-block`

```
#pragma omp parallel
{
    #pragma omp single
    {
        #pragma omp task
        fred();
        #pragma omp task
        daisy();
        #pragma taskwait // fred() and daisy() must complete before billy() starts
        #pragma omp task
        billy();
    } // All tasks complete before this barrier is released
}
```

# Data sharing: task defaults

- The behavior you want for tasks is usually **firstprivate**, because the task may not be executed until later (and variables may have gone out of scope)
  - Variables that are **private** when the task construct is encountered are firstprivate by default
- Variables that are shared in all constructs starting from the innermost enclosing parallel construct are shared by default

```
#pragma omp parallel shared(A) private(B)
```

```
{
```

```
...
```

```
#pragma omp task
```

```
{
```

```
    int C;
```

```
    compute(A, B, C); // B is firstprivate, C is private
```

```
}
```

```
}
```

# Serial Fibonacci Numbers

```
int fib (int n) {  
    int x,y;  
    if (n < 2) return n;  
  
    x = fib(n-1);  
    y = fib(n-2);  
    return (x+y);  
}
```

```
int main() {  
    int NW = 5000;  
    fib(NW);  
}
```

- $F_n = F_{n-1} + F_{n-2}$
- Inefficient  $O(2^n)$   
recursive  
implementation!



# Parallel Fibonacci Numbers

```
int fib (int n) {  int x,y;
    if (n < 2) return n;
    #pragma omp task shared(x)
        x = fib(n-1);
    #pragma omp task shared(y)
        y = fib(n-2);
    #pragma omp taskwait
        return (x+y);
}

int main() {  int NW = 5000;
    #pragma omp parallel
    {
        #pragma omp single
            fib(NW);
    }
}
```

- Binary tree of tasks
- Traversed using a recursive function
- A task cannot complete until all tasks below it in the tree are complete (enforced with **taskwait**)
- By default, only 2 threads will be active in most implementations. Set `OMP_MAX_ACTIVE_LEVELS` with `n>1` to get n-levels of nested parallelism
- x,y are local, and so by default they are private to current task
  - must be shared on child tasks so they don't create their own **firstprivate** copies at this level!

# Synchronization (slight return)

- High level synchronization:
  - critical
  - barrier
  - atomic
  - ordered
- Low level synchronization
  - flush
  - locks (both simple and nested)

Synchronization is used to impose order constraints and to protect access to shared data

# Synchronization: atomic

- Atomic provides mutual exclusion but only applies to the update of a memory location (the update of X in the following example)

```
#pragma omp parallel
```

```
{
```

```
    double B;
```

```
    B = DOIT();
```

```
#pragma omp atomic // faster than critical
```

```
    X += big_ugly(B);
```

```
}
```

# Low-level synchronization: flush

- Defines a sequence point at which a thread enforces a consistent view of memory.
- For variables visible to other threads and associated with the flush operation (the flush-set)
  - The compiler can't move loads/stores of the flush-set around a flush:
  - All previous read/writes of the flush-set by this thread have completed
  - No subsequent read/writes of the flush-set by this thread have occurred
  - Variables in the flush set are moved from temporary storage to shared memory.
  - Reads of variables in the flush set following the flush are loaded from shared memory.

# Low-level synchronization: flush

- A flush operation is implied by OpenMP synchronizations, e.g.,
  - at entry/exit of parallel regions
  - at implicit and explicit barriers
  - at entry/exit of critical regions
  - whenever a lock is set or unset

# Takeaway

- Programming shared memory machines
  - May allocate data in large shared region without too many worries about where
  - Memory hierarchy is critical to performance
    - Even more so than on uniprocessors, due to coherence traffic
  - For performance tuning, watch sharing (both true and false)
- Semantics
  - Need to lock access to shared variable for read-modify-write
  - Sequential consistency is the natural semantics
    - Write race-free programs to get this
  - Architects worked hard to make this work
    - Caches are coherent with buses or directories
    - No caching of remote data on shared address space machines
  - But compiler and processor may still get in the way
    - Non-blocking writes, read prefetching, code motion...
    - Avoid races or use machine-specific fences carefully

# Shared Memory Programming

- PTHREADS is the POSIX Standard
  - Portable but; relatively heavyweight; low level
- OpenMP standard for application level programming
  - Support for scientific programming on shared memory, [openmp.org](http://openmp.org)
- TBB: Thread Building Blocks
  - Intel C++ template library for parallel (multicore) computing
- Java threads
  - Built on top of POSIX threads; Object within Java language

# Parallel programming and Simulation

- Parallelism and data locality both critical to performance
  - Recall that moving data is the most expensive operation
- Real world problems have parallelism and locality:
  - Many objects operate independently of others.
  - Objects often depend much more on nearby than distant objects.
  - Dependence on distant objects can often be simplified.
  - Example of all three: particles moving under gravity
- Scientific models may introduce more parallelism:
  - When a continuous problem is discretized, time dependencies are generally limited to adjacent time steps.
    - Helps limit dependence to nearby objects (eg collisions)
  - Far-field effects may be ignored or approximated in many cases.
- Many problems exhibit parallelism at multiple levels



# Kinds of Simulation (from discrete to continuous)

- Discrete event systems:
  - “Game of Life,” Manufacturing systems, Finance, Circuits, Pacman, ...
- Particle systems:
  - Billiard balls, Galaxies, Atoms, Circuits, Pinball ...
- Lumped variables depending on continuous parameters: a.k.a. Ordinary Differential Equations (ODEs)
  - Structural mechanics, Chemical kinetics, Circuits, Star Wars: The Force Unleashed
- Continuous variables depending on continuous parameters: a.k.a. Partial Differential Equations (PDEs)
  - Heat, Elasticity, Electrostatics, Finance, Circuits, Medical Image Analysis, Terminator 3: Rise of the Machines

# Simulation outline

- Discrete event systems
  - Time and space are discrete
- Particle systems
  - Important special case of lumped systems
- Lumped systems (ODEs)
  - Location/entities are discrete, time is continuous
- Continuous systems (PDEs)
  - Time and space are continuous

# Circuit Simulation

Level	Primitives	Example
Instruction	Instructions	SPIM
Cycle	Functional Units	VHDL/ Verilog
RTL	Registers, muxes, etc	VHDL/ Verilog
Gate	Gate, FF, memory	VHDL/ Verilog
Switch	Transistor	COSMOS
Circuit	Resistor, Capacitor	SPICE
Device	Electrons, Semiconductor	...

# A Model Problem: Sharks and Fish

- Illustration of parallel programming
- Original version (discrete event only) called WATOR
- Basic idea: sharks and fish living in an ocean
  - rules for movement (discrete and continuous)
  - breeding, eating, and death
  - forces in the ocean
  - forces between sea creatures
- 6 problems (S&F1 - S&F6)
  - Different sets of rules, to illustrate different phenomena

# A Model Problem: Sharks and Fish

- S&F 1. Fish alone move continuously subject to an external current and Newton's laws (*completely parallel*)
- S&F 2. Fish alone move continuously subject to gravitational attraction and Newton's laws
- S&F 3. Fish alone play the "Game of Life" on a square grid
- S&F 4. Fish alone move randomly on a square grid, with at most one fish per grid point (*need to keep at most one fish per cell, so need to deal with possible races, deadlock*)
- S&F 5. Sharks and Fish both move randomly on a square grid, with at most one fish or shark per grid point, including rules for fish attracting sharks, eating, breeding and dying
- S&F 6. Like Sharks and Fish 5, but continuous, subject to Newton's laws

# Discrete Event Systems

- Systems are represented as:
  - finite set of variables.
  - the set of all variable values at a given time is called the state.
  - each variable is updated by computing a *transition function* depending on the other variables.
- System may be:
  - synchronous: at each discrete timestep evaluate all transition functions; also called a state machine.
  - asynchronous: transition functions are evaluated only if the inputs change, based on an “event” from another part of the system; also called event driven simulation.
- Example: The “game of life:” (Conway)
  - Also known as Sharks and Fish #3:
  - Space divided into cells, rules govern cell contents at each step

# Parallelism in Game of Life (S&F 3)

- The simulation is synchronous
  - use two copies of the grid (old and new), “ping-pong” between them
  - the value of each new grid cell depends only on 9 cells (itself plus 8 neighbors) in old grid.
  - simulation proceeds in timesteps-- each cell is updated at every step.
- Easy to parallelize by dividing physical domain: Domain Decomposition
  - Locality is achieved by using large patches of the ocean
- Only boundary values from neighboring patches are needed.
- How to pick shapes of domains?

# Game of Life (S&F 3)

- Only need two grids: ping pong between them, avoid race conditions
- Load balanced because each processor gets equal #grid cells
- Square domains seem natural, but lots of choices in general, how choose?
- Rules:
  - Any live cell with two or three live neighbors survives.
  - Any dead cell with three live neighbors becomes a live cell.
  - All other live cells die in the next generation. Similarly, all other dead cells stay dead.

Repeat

compute locally to update local system

barrier()

exchange state info with neighbors

finish updates

until done simulating



# Synchronous Circuit Simulation

- Circuit is a graph made up of subcircuits connected by wires
  - Component simulations need to interact if they share a wire.
  - Data structure is (irregular) graph of subcircuits.
  - Parallel algorithm is timing-driven or synchronous:
    - Evaluate all components at every *timestep* (determined by known circuit delay)
- Graph partitioning assigns subgraphs to processors
  - Determines parallelism and locality.
  - Goal 1 is to evenly distribute subgraphs to nodes (load balance).
  - Goal 2 is to minimize edge crossings (minimize communication).
  - Easy for meshes, NP-hard in general, so we will approximate (future lecture)

# Asynchronous Simulation

- Synchronous simulations may waste time:
  - Simulates even when the inputs do not change
- Asynchronous (event-driven) simulations update only when an event arrives from another component:
  - No global time steps, but individual events contain time stamp.
  - Example: Game of life in loosely connected ponds (don't simulate empty ponds).
  - Example: Circuit simulation with delays (events are gates changing).
  - Example: Traffic simulation (events are cars changing lanes, etc.).
- Asynchronous is more efficient, but harder to parallelize
  - On distributed memory, events are naturally implemented as messages between processors (eg using MPI), but how do you know when to execute a “receive”?

# Scheduling Asynchronous Circuit Simulation

- Conservative:
  - Only simulate up to (and including) the minimum time stamp of inputs.
  - Need deadlock detection if there are cycles in graph
  - Example: Pthor circuit simulator
- Speculative (or Optimistic):
  - Assume no new inputs will arrive and keep simulating.
  - May need to backup if assumption wrong, using timestamps
- Example: Timewarp [D. Jefferson], Parswec [Wen, Yelick].
- Optimizing load balance and locality is difficult:
  - Locality means putting tightly coupled subcircuit on one processor.
  - Since “active” part of circuit likely to be in a tightly coupled subcircuit, this may be bad for load balance.

# Summary of Discrete Event Simulations

- Model of the world is discrete
  - Both time and space
- Approaches
  - Decompose domain, i.e., set of objects
  - Run each component ahead using
    - Synchronous: communicate at end of each timestep
    - Asynchronous: communicate on-demand
      - Conservative scheduling – wait for inputs
        - need deadlock detection
      - Speculative scheduling – assume no inputs
        - roll back if necessary