

Outline

- Incredibly brief review of simulation
- On to Differential Equations: ODEs and PDEs
- A brief look at tree algorithms (and Scan)

Simulation outline

- Discrete event systems
 - Time and space are discrete
- Particle systems
 - Important special case of lumped systems
- Lumped systems (ODEs)
 - Location/entities are discrete, time is continuous
- Continuous systems (PDEs)
 - Time and space are continuous

Summary of particle methods

- Model contains discrete entities, namely, particles
- Time is continuous – must be discretized to solve
- Simulation follows particles through timesteps
 - $\text{Force} = \text{external_force} + \text{nearby_force} + \text{far_field_force}$
 - All-pairs algorithm is simple, but inefficient, $O(n^2)$
 - *Particle-mesh* methods approximate by moving particles to a regular mesh, where it is easier to compute forces
 - *Tree-based* algorithms approximate by treating set of particles as a group, when far away

Review of last lecture

- Common problems:
 - Load balancing
 - May be due to lack of parallelism or poor work distribution
 - Statically, divide grid (or graph) into blocks
 - Dynamically, if load changes significantly during run
 - Locality
 - Partition into large chunks with low surface-to-volume ratio
 - To minimize communication
 - Distributed particles according to location, but use irregular spatial decomposition (e.g., quad tree) for load balance
 - Constant tension between these two
 - Particle-Mesh method: can't balance particles (moving), balance mesh (fixed) and keep particles near mesh points without communication

System of Lumped Variables

- Many systems are approximated by
 - System of “lumped” variables.
 - Each depends on continuous parameter (usually time).
- Example -- circuit:
 - approximate as graph.
 - wires are edges.
 - nodes are connections between 2 or more wires.
 - each edge has resistor, capacitor, inductor or voltage source.
 - system is “lumped” because we are not computing the voltage/current at every point in space along a wire, just endpoints.
 - Variables related by Ohm’s Law, Kirchhoff’s Laws, etc.
- Forms a system of *ordinary differential equations* (ODEs).
 - Differentiated with respect to time
 - Variant: ODEs with some constraints

Circuit Example

- State of the system is represented by
 - $v_n(t)$ node voltages
 - $i_b(t)$ branch currents all at time t
 - $v_b(t)$ branch voltages

- Equations include

- Kirchoff's current
- Kirchoff's voltage
- Ohm's law
- Capacitance
- Inductance

$$\begin{pmatrix} 0 & A & 0 \\ A' & 0 & -I \\ 0 & R & -I \\ 0 & -I & C*d/dt \\ 0 & L*d/dt & I \end{pmatrix} * \begin{pmatrix} v_n \\ i_b \\ v_b \end{pmatrix} = \begin{pmatrix} 0 \\ S \\ 0 \\ 0 \\ 0 \end{pmatrix}$$

- A is sparse matrix, representing connections in circuit
 - One column per branch (edge), one row per node (vertex) with +1 and -1 in each column at rows indicating end points
- Write as single large system of ODEs

Solving ODEs

- In these examples, and most others, the matrices are sparse:
 - i.e., most array elements are 0.
 - neither store nor compute on these 0's.
 - Sparse because each component only depends on a few others
- Given a set of ODEs, two kinds of questions are:
 - Compute the values of the variables at some time t
 - Explicit methods
 - Implicit methods
 - Compute modes of vibration
 - Eigenvalue problems

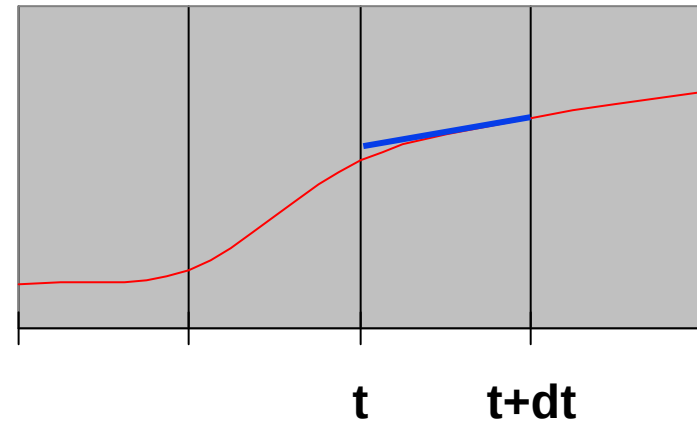

Solving ODEs: Explicit methods

- Assume ODE is $x'(t) = f(x) = A*x(t)$, where A is a sparse matrix
 - Compute $x(i*dt) = x[i]$ at $i=0,1,2,\dots$
 - ODE gives $x'(i*dt) = \text{slope}$
$$x[i+1] = x[i] + dt * \text{slope}$$
- Explicit methods, e.g., (Forward) Euler's method.
 - Approximate $x'(t) = A*x(t)$ by $(x[i+1] - x[i]) / dt = A*x[i]$.
 - $x[i+1] = x[i] + dt * A*x[i]$, i.e. sparse matrix-vector multiplication.
- Tradeoffs:
 - Simple algorithm: sparse matrix vector multiply.
 - Stability problems: May need to take very small time steps, especially if system is “stiff” (i.e. A has some large entries, so x can change rapidly).

Solving ODEs: Implicit methods

- Assume ODE is $x'(t) = f(x) = A*x(t)$, where A is a sparse matrix
 - Compute $x(i*dt) = x[i]$ at $i=0,1,2,\dots$
 - ODE gives $x'((i+1)*dt) = \text{slope}$
 $x[i+1]=x[i] + dt*\text{slope}$
- Implicit method, e.g., Backward Euler solve:
 - Approximate $x'(t)=A*x(t)$ by $(x[i+1] - x[i])/dt = A*x[i+1]$.
 - $(I - dt*A)*x[i+1] = x[i]$, i.e. we need to solve a sparse linear system of equations.
- Trade-offs:
 - Larger timestep possible: especially for stiff problems
 - More difficult algorithm: need to solve a sparse linear system of equations at each step

Explicit vs. Implicit



Forward (Explicit: $dt(i+1)$) vs. Backward (Implicit): dt

Solving ODEs: Eigenvalue methods

- Computing modes of vibration: finding eigenvalues and eigenvectors.
 - Seek solution of $d^2 x(t)/dt^2 = A*x(t)$ of form
$$x(t) = \sin(\omega*t) * x_0, \text{ where } x_0 \text{ is a constant vector}$$
 - ω called the frequency of vibration
 - x_0 sometimes called a “mode shape”
 - Plug in to get $-\omega^2 * x_0 = A*x_0$, so that $-\omega^2$ is an eigenvalue and x_0 is an eigenvector of A .
 - Solution schemes reduce either to sparse-matrix multiplication, or solving sparse linear systems.

Implicit methods: Eigenproblems

- Implicit methods for ODEs need to solve linear systems
- Direct methods (Gaussian elimination)
 - Called LU Decomposition, because we factor $A = L*U$
 - More complicated than sparse-matrix vector multiplication.
- Iterative solvers
 - Jacobi, Successive over-relaxation (SOR) , Conjugate Gradient (CG), Multigrid,...
 - Most have sparse-matrix-vector multiplication
- Eigenproblems
 - Also depend on sparse-matrix-vector multiplication, direct methods.

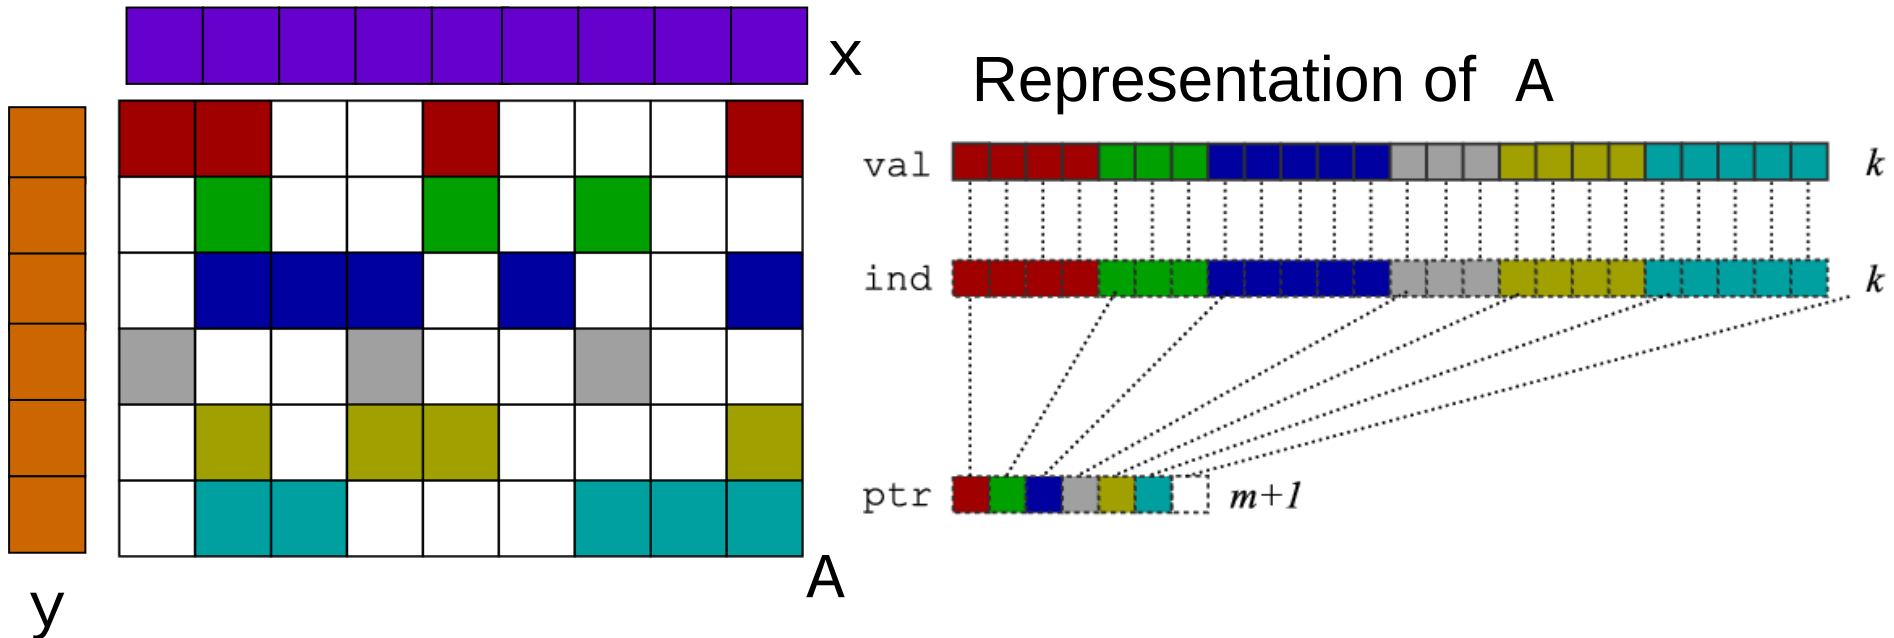
ODEs and Sparse Matrices

- All these problems reduce to sparse matrix problems
 - Explicit: sparse matrix-vector multiplication
 - Implicit: solve a sparse linear system
 - direct solvers (Gaussian elimination).
 - iterative solvers (use sparse matrix-vector multiplication).
 - Eigenvalue/vector algorithms may also be explicit or implicit.
- Conclusion: *Sparse Matrix-Vector Multiplication is key* to many ODE problems
 - Relatively simple algorithm to study in detail
 - Two key problems: locality and load balance

Compressed Sparse Row (CSR) Format

$y = y + A \cdot x$, only store, do arithmetic, on nonzero entries

CSR format is simplest one of many possible data structures for A



Matrix-vector multiply kernel: $y(i) \leftarrow y(i) + A(i,j) \cdot x(j)$


for each row i

for $k=ptr[i]$ to $ptr[i+1]-1$ do

$y[i] = y[i] + val[k] \cdot x[ind[k]]$

Parallel Sparse Matrix-vector multiplication

- $y = A * x$, where A is a sparse $n \times n$ matrix




- Questions
 - which processors store
 - $y[i]$, $x[i]$, and $A[i,j]$
 - which processors compute
 - $y[i] = \text{sum (from 1 to } n) A[i,j] * x[j]$
 $= (\text{row } i \text{ of } A) * x$... a sparse dot product
- Partitioning
 - Partition index set $\{1, \dots, n\} = N1 \cup N2 \cup \dots \cup Np$.
 - For all i in Nk , Processor k stores $y[i]$, $x[i]$, and row i of A
 - For all i in Nk , Processor k computes $y[i] = (\text{row } i \text{ of } A) * x$
 - “owner computes” rule: Processor k compute the $y[i]$ s it owns.

May require communication

Matrix Reordering via Graph Partitioning

- “Ideal” matrix structure for parallelism: block diagonal
 - p (number of processors) blocks, can all be computed locally.
 - If no non-zeros outside these blocks, no communication needed
- Can we reorder the rows/columns to get close to this
 - Most nonzeros in diagonal blocks, few outside




Goals of Reordering

- Performance goals
 - balance load (how is load measured?).
 - Approx equal number of nonzeros (not necessarily rows)
 - balance storage (how much does each processor store?).
 - Approx equal number of nonzeros
 - minimize communication (how much is communicated?).
 - Minimize nonzeros outside diagonal blocks
 - Related optimization criterion is to move nonzeros near diagonal
 - improve register and cache re-use
 - Group nonzeros in small vertical blocks so source (x) elements loaded into cache or registers may be reused (temporal locality)
 - Group nonzeros in small horizontal blocks so nearby source (x) elements in the cache may be used (spatial locality)
- Other algorithms reorder for other reasons
 - Reduce # nonzeros in matrix after Gaussian elimination
 - Improve numerical stability

Graph Partitioning

- Relationship between matrix and graph

	1	2	3	4	5	6
1	1			1	1	
2		1	1	1	1	
3		1	1			1
4	1	1		1		1
5	1	1			1	1
6			1	1	1	1




- Edges in the graph are nonzero in the matrix: here the matrix is symmetric (edges are unordered) and weights are equal (1)
- If divided over 3 processors, there are 14 nonzeros outside the diagonal blocks, which represent the 7 (bidirectional) edges

Graph Partitioning and Sparse Matrices

- Relationship between matrix and graph

	1	2	3	4	5	6
1	1	1			1	
2	1	1		1		1
3			1	1		1
4		1	1	1	1	
5	1			1	1	1
6		1	1		1	1



- A “good” partition of the graph has
 - equal (weighted) number of nodes in each part (load and storage balance).
 - minimum number of edges crossing between (minimize communication).
- Reorder the rows/columns by putting all nodes in one partition together.

Summary of common problems

- Load Balancing
 - Dynamically – if load changes significantly during job
 - Statically - Graph partitioning
 - Discrete systems
 - Sparse matrix vector multiplication
- Linear algebra
 - Solving linear systems (sparse and dense)
 - Eigenvalue problems will use similar techniques
- Fast Particle Methods
 - $O(n \log n)$ instead of $O(n^2)$

Computational methods in Applications

	Embed	SPEC	DB	Games	ML	HPC	Health	Image	Speech	Music	Browser
1 Finite State Mach.	Red	Red	Red	Yellow	Yellow	LightBlue	LightBlue	LightBlue	LightBlue	LightBlue	Red
2 Combinational	Red	LightBlue	Green	LightBlue	Green	LightBlue	LightBlue	LightBlue	LightBlue	LightBlue	Red
3 Graph Traversal	Red	Yellow	Yellow	Yellow	Red	LightBlue	Red	LightBlue	Red	Green	Green
4 Structured Grid	Red	Red	LightBlue	Yellow	LightBlue	Red	LightBlue	Red	LightBlue	LightBlue	LightBlue
5 Dense Matrix	Red	Red	Yellow	Red	Red	Red	LightBlue	Red	Red	Red	LightBlue
6 Sparse Matrix	Yellow	Yellow	LightBlue	Red	Red	Red	Red	LightBlue	LightBlue	Red	LightBlue
7 Spectral (FFT)	Yellow	LightBlue	LightBlue	Yellow	Yellow	Red	LightBlue	Green	Red	Red	Red
8 Dynamic Prog	Yellow	LightBlue	Red	LightBlue	Red	LightBlue	LightBlue	LightBlue	Yellow	LightBlue	Red
9 N-Body	LightBlue	Yellow	LightBlue	Yellow	LightBlue	Red	Green	LightBlue	LightBlue	LightBlue	LightBlue
10 MapReduce	LightBlue	Green	Red	LightBlue	Red	Red	Red	Red	Yellow	Red	Yellow
11 Backtrack/ B&B	LightBlue	LightBlue	Yellow	LightBlue	Red	LightBlue	LightBlue	LightBlue	LightBlue	Yellow	LightBlue
12 Graphical Models	LightBlue	LightBlue	Yellow	LightBlue	Red	LightBlue	LightBlue	LightBlue	LightBlue	Red	LightBlue
13 Unstructured Grid	LightBlue	LightBlue	LightBlue	Yellow	Yellow	Red	Red	LightBlue	LightBlue	Red	LightBlue

PDEs: Continuous Variables, Continuous Parameters

- Examples of such systems include
 - Elliptic problems (steady state, global space dependence)
 - Electrostatic or Gravitational Potential: $\text{Potential}(\text{position})$
 - Hyperbolic problems (time dependent, local space dependence):
 - Sound waves: $\text{Pressure}(\text{position}, \text{time})$
 - Parabolic problems (time dependent, global space dependence)
 - Heat flow: $\text{Temperature}(\text{position}, \text{time})$
 - Diffusion: $\text{Concentration}(\text{position}, \text{time})$

PDEs: Local/Global Dependence

- Global vs Local Dependence
 - Global means either a lot of communication, or tiny time steps
 - Local arises from finite wave speeds: limits communication
- Many problems combine features of above
 - Fluid flow: Velocity, Pressure, Density(position, time)
 - Elasticity: Stress, Strain(position, time)

Explicit time stepping

- Approximate PDE by ODE system (“method of lines”):
- Need a time-stepping scheme for the ODE: Simplest scheme is Euler’s Method
- Taking a time step \equiv sparse matrix vector multiplication
- This may not end well... (instability)

Implicit time stepping

- Examples of such systems include
 - Elliptic problems (steady state, global space dependence)
 - Electrostatic or Gravitational Potential: $\text{Potential}(\text{position})$
 - Hyperbolic problems (time dependent, local space dependence):
 - Sound waves: $\text{Pressure}(\text{position}, \text{time})$
 - Parabolic problems (time dependent, global space dependence)
 - Heat flow: $\text{Temperature}(\text{position}, \text{time})$
 - Diffusion: $\text{Concentration}(\text{position}, \text{time})$

Parallelism in Explicit Method for PDEs

- Sparse matrix vector multiply, via Graph Partitioning
- Partitioning the space (x) into p chunks
 - good load balance (assuming large number of points relative to p)
 - minimize communication (least dependence on data outside chunk)
- Generalizes to
 - multiple dimensions.
 - arbitrary graphs (= arbitrary sparse matrices).
- Explicit approach often used for hyperbolic equations
 - Finite wave speed, so only depend on nearest chunks
- Problem with explicit approach for heat (parabolic): numerical instability.

Implicit vs. Explicit

- Explicit:
 - Propagates information at finite rate
 - Steps look like sparse matrix-vector (in linear case)
 - Stable step determined by fastest time scale
 - Works fine for hyperbolic PDEs
- Implicit:
 - No need to resolve fastest time scales
 - Steps can be long... but expensive
 - Linear/nonlinear solves at each step
 - Often these solves involve sparse matrix-vectors
 - Critical for parabolic PDEs

Algorithm overview

from slowest to fastest on sequential machines

- **Dense LU**: Gaussian elimination; works on any N -by- N matrix.
- **Band LU**: Exploits the fact that tridiagonal matrix T is nonzero only on \sqrt{N} diagonals nearest main diagonal.
- **Jacobi**: Essentially does matrix-vector multiply by T in inner loop of iterative algorithm.
- **Explicit Inverse**: Assume we want to solve many systems with T , so we can precompute and store $\text{inv}(T)$ “for free”, and just multiply by it (but still expensive).
- **Conjugate Gradient**: Uses matrix-vector multiplication, like Jacobi, but exploits mathematical properties of T that Jacobi does not.
- **Red-Black SOR (successive over-relaxation)**: Variation of Jacobi that exploits yet different mathematical properties of T . Used in multigrid schemes.
- **Sparse LU**: Gaussian elimination exploiting particular zero structure of T .
- **FFT (Fast Fourier Transform)**: Works only on matrices very like T .
- **Multigrid**: Also works on matrices like T , that come from elliptic PDEs.
- **Lower Bound**: Serial (time to print answer); parallel (time to combine N inputs).

Summary of Approaches to Solving PDEs


- As with ODEs, either explicit or implicit approaches are possible
 - Explicit, sparse matrix-vector multiplication
 - Implicit, sparse matrix solve at each step
 - Direct solvers are hard
 - Iterative solves turn into sparse matrix-vector multiplication: Graph partitioning
- Graph and sparse matrix correspondence:
 - Sparse matrix-vector multiplication is nearest neighbor “averaging” on the underlying mesh
- Not all nearest neighbor computations have the same efficiency
 - Depends on the mesh structure (nonzero structure) and the number of Flops per point.

Comments on practical meshes

- Regular 1D, 2D, 3D meshes
 - Important as building blocks for more complicated meshes
- Practical meshes are often irregular
 - **Composite meshes**, consisting of multiple “bent” regular meshes joined at edges
 - **Unstructured meshes**, with arbitrary mesh points and connectivities
 - **Adaptive meshes**, which change resolution during solution process to put computational effort where needed

Parallelism in Regular meshes


- Computing a Stencil on a regular mesh
 - need to communicate mesh points near boundary to neighboring processors.
 - Often done with ghost regions
- Surface-to-volume ratio keeps communication down, but
 - Still may be problematic in practice




Implemented using “ghost” regions.

Adds memory overhead

Irregular mesh: NASA Airfoil in 2D (direct solution)



Adaptive mesh



Challenges of Irregular Meshes


- How to generate them in the first place
 - Start from geometric description of object
 - Triangle, a 2D mesh partitioner by Jonathan Shewchuk
 - 3D harder!
- How to partition them
 - ParMetis, a parallel graph partitioner
- How to design iterative solvers
 - PETSc, a Portable Extensible Toolkit for Scientific Computing
 - Prometheus, a multigrid solver for finite element problems on irregular meshes
- How to design direct solvers
 - SuperLU, parallel sparse Gaussian elimination

The “Seven Dwarfs”:

High-end simulation in the physical sciences

- 1) Structured grids
- 2) Unstructured grids
- 3) Spectral methods (Fast Fourier Transform)
- 4) Dense Linear Algebra
- 5) Sparse Linear Algebra: Both explicit and implicit
- 6) Particle Methods
- 7) Monte Carlo/Embarrassing Parallelism/Map Reduce
(easy!)

Tree structured computation




Parallel Prefix, or Scan

- If “+” is an associative operator, and $x[0], \dots, x[p-1]$ are input data then parallel prefix operation computes

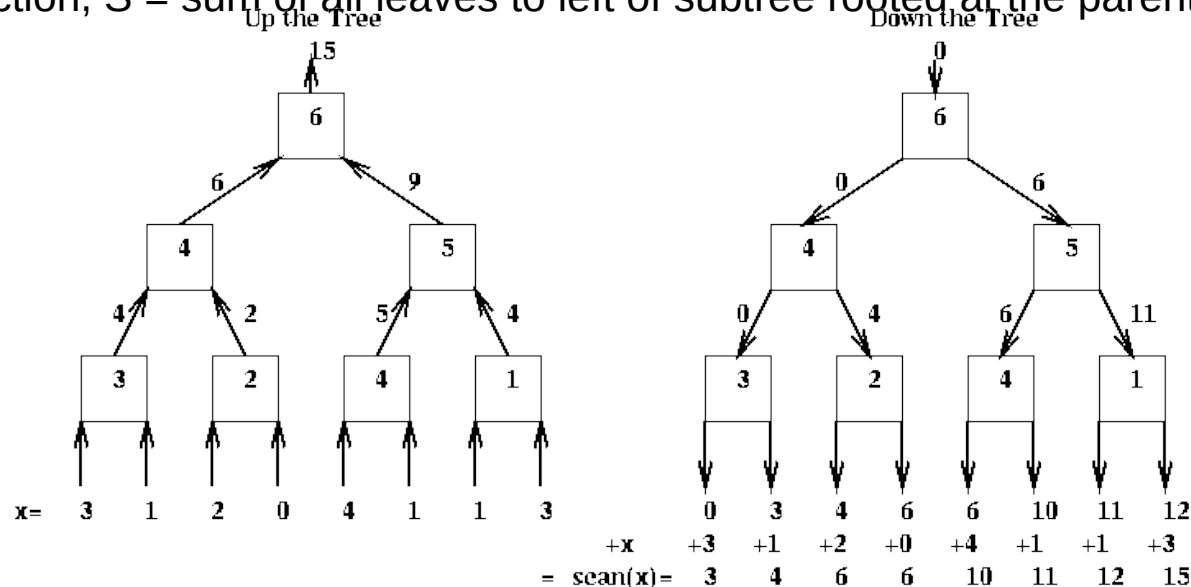
$$y[j] = x[0] + x[1] + \dots + x[j] \quad \text{for } j=0,1,\dots,p-1$$

- Notation: $j:k$ means $x[j]+x[j+1]+\dots+x[k]$, **blue** is final value



Mapping Parallel Prefix onto a Tree: Details

- Up-the-tree phase (from leaves to root)
 - 1) Get values L and R from left and right children
 - 2) Save L in a local register Lsave
 - 3) Pass sum L+R to parent
- By induction, Lsave = sum of all leaves in left subtree
- Down the tree phase (from root to leaves)
 - 1) Get value S from parent (the root gets 0)
 - 2) Send S to the left child
 - 3) Send S + Lsave to the right child
- By induction, S = sum of all leaves to left of subtree rooted at the parent



Adding two n-bit ints in $O(\log n)$ time

- Let $a = a[n-1]a[n-2]\dots a[0]$ and $b = b[n-1]b[n-2]\dots b[0]$ be two n-bit binary numbers

- We want their sum $s = a+b = s[n]s[n-1]\dots s[0]$

$c[-1] = 0$... rightmost carry bit

for $i = 0$ to $n-1$

$c[i] = ((a[i] \text{ xor } b[i]) \text{ and } c[i-1]) \text{ or } (a[i] \text{ and } b[i])$... next carry bit

$s[i] = (a[i] \text{ xor } b[i]) \text{ xor } c[i-1]$

- Challenge: compute all $c[i]$ in $O(\log n)$ time via parallel prefix
 for all $(0 \leq i \leq n-1)$ $p[i] = a[i] \text{ xor } b[i]$... propagate bit
 for all $(0 \leq i \leq n-1)$ $g[i] = a[i] \text{ and } b[i]$... generate bit

$$\begin{bmatrix} c[i] \\ 1 \end{bmatrix} = \begin{bmatrix} (p[i] \text{ and } c[i-1]) \text{ or } g[i] \\ 1 \end{bmatrix} = \begin{bmatrix} p[i] & g[i] \\ 0 & 1 \end{bmatrix} * \begin{bmatrix} c[i-1] \\ 1 \end{bmatrix} = C[i] * \begin{bmatrix} c[i-1] \\ 1 \end{bmatrix}$$

... 2-by-2 Boolean matrix multiplication (associative)

$$= C[i] * C[i-1] * \dots * C[0] * \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

... evaluate each $P[i] = C[i] * C[i-1] * \dots * C[0]$ by parallel prefix

Used in all computers to implement addition - Carry look-ahead

Browser page layout via Prefix (Scan)

- Applying layout rules to html description of a webpage is a bottleneck, scan can help
- Simplest example
 - Given widths $[x_1, x_2, \dots, x_n]$ of items to display on page, where should each item go?
 - Item j starts at $x_1 + x_2 + \dots + x_{j-1}$
- Real examples have complicated constraints
 - Defined by general trees, since in html each object to display can be composed of other objects
 - To get location of each object, need to do preorder traversal of tree, “adding up” constraints of previous objects
 - Scan can do preorder traversal of any tree in parallel
 - Not just binary trees

Summary of tree algorithms

- Lots of problems can be done quickly - in theory - using trees
- Some algorithms are widely used
 - broadcasts, reductions, parallel prefix
 - carry look ahead addition
- Some are of theoretical interest only
 - Csanky's method for matrix inversion
 - Solving tridiagonal linear systems (without pivoting)
 - Both numerically unstable
- Embedded in various systems
 - MPI, NESL (CMU), other languages
 - CM-5 hardware control network