

Lecture plan

- ▶ Last time: Homework, Sparse Linear Algebra
- ▶ This time: FFT (and more)

Preliminaries: Laplace/Fourier Transform

- ▶ Laplace Transform $F(s) = \int_0^{\infty} f(t)e^{-st}dt$
 - ▶ $s = j\omega$, ω *continuous*
 - ▶ t is also *continuous*
- ▶ Fourier Transform $F(\omega) = \int_{-\infty}^{\infty} f(t)e^{-j\omega t}dt$

Continuous vs. Discrete

- ▶ But computers are **discrete!**
 - ▶ Both floating point and ints are fixed resolution
 - ▶ Time must be discretized (memory has discrete addresses)
- ▶ Solutions:
 - ▶ *Quantize* continuous inputs → discrete (integer) values
 - ▶ *Sampling* continuous time → discrete time

DFT definitions

- ▶ Define $j = \sqrt{-1}$ and index matrices and vectors from 0.
- ▶ Given some number N , let
$$W_N = e^{-j2\pi/N} = \cos(2\pi/N) + j * \sin(2\pi/N)$$
- ▶ W is a complex number with whose N -th power $W_N^N = 1$ and is therefore called an N -th *root of unity*
- ▶ Used in the FFT, these powers of W are also called *twiddle factors*

Motivation for the FFT

- ▶ Signal processing
- ▶ Image processing
- ▶ Solving Poisson's Equation nearly optimally
 - ▶ $O(N \log N)$ arithmetic operations, N = number of unknowns
 - ▶ Competitive with multigrid
- ▶ Fast multiplication of large integers: Schonhage-Strassen:
 $O(b \log b \log \log b)$ where b = number of bits

Poisson's equation arises in many models

In general, $\sum_1^N \frac{\partial^2 \phi}{\partial x_i^2} = f$

- ▶ 3D: $\partial^2 u / \partial x^2 + \partial^2 u / \partial y^2 + \partial^2 u / \partial z^2 = f(x, y, z)$
- ▶ 2D: $\partial^2 u / \partial x^2 + \partial^2 u / \partial y^2 = f(x, y)$
- ▶ 1D: $d^2 u / dx^2 = f(x)$

Applications:

- ▶ Electrostatic or Gravitational Potential: Potential(position)
- ▶ Heat flow: Temperature(position, time)
- ▶ Diffusion: Concentration(position, time)
- ▶ Fluid flow: Velocity, Pressure, Density(position, time)
- ▶ Elasticity: Stress, Strain(position, time)
- ▶ Variations of Poisson have variable coefficients

2D Poisson with FFT

- ▶ $L_1 = F \cdot D \cdot F^T$ is eigenvalue/eigenvector decomposition, where
 - ▶ F is very similar to FFT (imaginary part):
$$F(j, k) = (2/(n+1))^{1/2} \cdot \sin(jkp/(n+1))$$
 - ▶ D = diagonal matrix of eigenvalues:
$$D(j, j) = 2(1 - \cos(jp/(n+1)))$$
- ▶ 2D Poisson same as solving $L_1 \cdot X + X \cdot L_1 = B$ where X square matrix of unknowns at each grid point, B square too
- ▶ Substitute $L_1 = F \cdot D \cdot F^T$ into 2D Poisson to get algorithm:
 1. Perform 2D “FFT” on B to get $B' = F^T \cdot B \cdot F$, or
$$B = F \cdot B' \cdot F^T$$
 2. Solve $DX' + X'D = B'$ for
$$X' : X'(j, k) = B'(j, k)/(D(j, j) + D(k, k))$$
 3. Perform inverse 2D “FFT” on $X' = F^T \cdot X \cdot F$ to get
$$X = F \cdot X' \cdot F^T$$
- ▶ Cost = 2 2D-FFTs plus n^2 adds, divisions = $O(n^2 \log n)$
- ▶ 3D Poisson analogous

Related Transforms

- ▶ Most applications require multiplication by both F and F^{-1}
- ▶ Multiplying by F and F^{-1} are essentially the same.
- ▶ $F^{-1} = F^*/m$, $*$ is the complex conjugate
- ▶ For solving the Poisson equation and various other applications, we use variations on the FFT
 - ▶ The sin transform – imaginary part of F
 - ▶ The cos transform – real part of F
- ▶ Algorithms are similar, so we will focus on F

Serial Algorithm for the FFT

Compute the $FFT(F * v)$ of an N -element vector v

$$\begin{aligned}(F * v)[j] &= \sum_{k=0}^{N-1} F(j, k) * v(k) \\&= \sum_{k=0}^{N-1} W_N^{j * k} * v(k) \\&= \sum_{k=0}^{N-1} (W_N^j)^k * v(k) \\&= V(W_N^j)\end{aligned}$$

where V is defined as the polynomial $V(x) = \sum_{k=0}^{N-1} x^k * v(k)$
So, the FFT is the *same* as evaluating a polynomial $V(x)$ with degree $N - 1$ at N different points

Divide and Conquer FFT

V can be evaluated using divide-and-conquer:

$$V(x) = \sum_{k=0}^{N-1} x^k * v(k) \quad (1)$$

$$= v[0] + x^2 * v[2] + x^4 * v[4] + \dots + \quad (2)$$

$$x * (v[1] + x^2 * v[3] + x^4 * v[5] + \dots) \quad (3)$$

$$= V_{\text{even}}(x^2) + x * V_{\text{odd}}(x^2) \quad (4)$$

$$(5)$$

- ▶ V has degree $N - 1$, so V_{even} and V_{odd} are polynomials of degree $N/2 - 1$
- ▶ We evaluate these at N points: $(W_N^l)^2$ for $0 \leq l \leq m - 1$
- ▶ But this is really just $N/2$ different points, since $(W^{(l + N/2)})^2 = (W^l * W^{N/2})^2 = W^{2l} * W^N = (W^l)^2$
- ▶ So FFT on N points reduced to 2 FFTs on $N/2$ points

Divide and conquer!

Divide and Conquer FFT

```
FFT(v, W, N) ... assume N is a power of 2
    if N = 1 return v[0]
    else
        veven = FFT(v[0:2:N-2],  $W_N^2$ , N/2)
        vodd = FFT(v[1:2:N-1],  $W_N^2$ , N/2)
        Wvec = [ $W_N^0, W_N^1, \dots, W_N^{(N/2-1)}$ ]
        return [veven + (Wvec .* vodd),
                veven - (Wvec .* vodd) ]
```

Cost: $T(N) = 2T(N/2) + O(N) = O(N \log N)$ operations

Twiddle factors are *precomputed*


Examining the call tree

- ▶ Divides by even and odd: low order bit
 1. level 0: xxx0, xxx1
 2. level 1: xx00, xx10 — xx01, xx11
 3. etc... etc...
- ▶ This is *bit reversed*!
- ▶ An iterative algorithm that uses loops rather than recursion, does each level in the tree starting at the bottom
- ▶ Algorithm overwrites $v[i]$ by $(F*v)[\text{bitreverse}(i)]$
- ▶ Practical algorithms combine recursion (for memory hierarchy) and iteration

Bit reversal


- ▶ The outputs of the FFT stored in *bit reversed* order
- ▶ May need to sort them (which is like transposes at various levels)
- ▶ Many FFT uses will do a reverse FFT as well, and therefore not require a bit reversal at all
- ▶ DSP processors offer bit reversed addressing

FFT by stage



FFT by stage: 4 processor mapping


Block Data Layout of an $m=16$ -point FFT on $p=4$ Processors



- ▶ Using a block layout (N/p contiguous words per processor)
- ▶ No communication in last $\log N/p$ steps
- ▶ Significant communication in first $\log p$ steps

FFT by stage: cyclic mapping

Cyclic Data Layout of an $m=16$ -point FFT on $p=4$ Processors



- ▶ Cyclic layout (consecutive words map to consecutive processors)
- ▶ No communication in first $\log(N/p)$ steps
- ▶ Communication in last $\log(p)$ steps


Parallel Complexity

- ▶ N = vector size, p = number of processors
- ▶ f = time per flop = 1
- ▶ a = latency for message
- ▶ b = time per word in a message

Time(blockFFT) = Time(cyclicFFT) =
 $2 * N * \log(N)/p$... perfectly parallel flops
+ $\log(p) * a$... 1 message/stage, $\log p$ stages
+ $N * \log(p)/p * b$... N/p words/message

FFT by stage: transpose mapping

Transpose Algorithm for an $m=16$ -point FFT on $p=4$ Processors



- ▶ Start with a cyclic layout for first $\log(N/p)$ steps, there is no communication
- ▶ Then transpose the vector for last $\log(p)$ steps
- ▶ All communication is in the transpose

Communication analogous to transposing an array

Sequential Communication Complexity of the FFT

How many words need to be moved between main memory and cache of size M to do the FFT of size N , where $N > M$?

- ▶ Theorem (Hong, Kung, 1981): words = $\Omega(N \log N / \log M)$:
Proof follows from each word of data being reusable only $\log M$ times; assumes no recomputation
- ▶ Attained by transpose algorithm
 - ▶ Sequential algorithm “simulates” parallel algorithm
 - ▶ Imagine we have $P = N/M$ processors, so each processor stores and works on $O(M)$ words
 - ▶ Each local computation phase in parallel FFT replaced by similar phase working on cache resident data in sequential FFT
 - ▶ Each communication phase in parallel FFT replaced by reading/writing data from/to cache in sequential FFT
- ▶ Attained by recursive, “cache-oblivious” algorithm (FFTW)

Higher Dimensional FFTs

- ▶ FFTs on two or more dimensions are defined as 1D FFTs on vectors in all dimensions.
- ▶ 2D FFT does 1D FFTs on all rows and then all columns
- ▶ There are 3 obvious possibilities for the 2D FFT:
 1. 2D blocked layout for matrix, using parallel 1D FFTs for each row and column
 2. Block row layout for matrix, using serial 1D FFTs on rows, followed by a transpose, then more serial 1D FFTs
 3. Block row layout for matrix, using serial 1D FFTs on rows, followed by parallel 1D FFTs on columns
- ▶ Option 2 is best, if we overlap communication and computation
- ▶ For a 3D FFT the options are similar
 - ▶ Two phases done with serial FFTs, followed by a transpose for 3rd
 - ▶ Can overlap communication with 2nd phase in practice

Bisection Bandwidth

- ▶ FFT requires one (or more) transpose operations: Every processor sends $1/p$ -th of its data to each other one
- ▶ Bisection Bandwidth limits this performance
 - ▶ Bisection bandwidth is the bandwidth across the narrowest part of the network
 - ▶ Important in global transpose operations, all-to-all, etc.
- ▶ “Full bisection bandwidth” is expensive
- ▶ Fraction of machine cost in the network is increasing
 - ▶ Fat-tree and full crossbar topologies may be too expensive
 - ▶ Especially on machines with 100K and more processors
 - ▶ SMP clusters often limit bandwidth at the node level
- ▶ Goal: overlap communication and computation

Fourier Transform Benchmark Case Study

- ▶ Performance of Exchange (All-to-all) is critical
 - ▶ Communication to computation ratio increases with faster, more optimized 1-D FFTs (used best available, from FFTW)
 - ▶ Determined by available bisection bandwidth
 - ▶ Between 30-40% of the application's total runtime
- ▶ Assumptions
 - ▶ 1D partition of 3D grid
 - ▶ At most N processors for N^3 grid
 - ▶ HPC Challenge benchmark has large 1D FFT (can be viewed as 3D or more with proper roots of unity)

3D FFTs

- ▶ $NX \times NY \times NZ$ elements spread across P processors
- ▶ Each processor gets NZ/P “planes” of $NX \times NY$ elements per plane
- ▶ 3 step process:
 1. FFTs on the columns (all elements local)
 2. FFTs on the rows (all elements local)
 3. FFTs in the Z-dimension (requires communication)

Parallel Processing

- ▶ Each processor has to scatter input domain to other processors
 - ▶ Every processor divides its portion of the domain into P pieces
 - ▶ Send each of the P pieces to a different processor
- ▶ Three different ways to break up the messages:
 1. Chunk (i.e. single packed “All-to-all” in MPI parlance) (3D) = all rows with same destination
 2. Slabs (2D) = all rows in a single plane with same destination
 3. Pencils (1D) = 1 row
- ▶ Going from Chunks to Slabs to Pencils leads to
 - ▶ An order of magnitude increase in the number of messages
 - ▶ An order of magnitude decrease in the size of each message
- ▶ Why do this? Slabs and Pencils allow *overlapping* communication and computation

Decomposing NAS FT Exchange into Smaller Messages

Three approaches:

1. Chunk: Wait for 2D FFTs to finish
2. Slab: Wait for chunk of rows destined for 1 proc to finish
3. Pencil: Send each row as it completes

Example Message Size Breakdown for Class D ($2048 \times 1024 \times 1024$) with 256 processors

Chunk 512 Kbytes

Slabs 65 Kbytes

Pencils 16 Kbytes

FFTW: Fastest Fourier Transform in the West

www.fftw.org

- ▶ Produces FFT implementation optimized for
 - ▶ Your version of FFT (complex, real,...)
 - ▶ Your value of N (arbitrary, possibly prime)
 - ▶ Your architecture
 - ▶ Very good sequential performance
- ▶ Won 1999 Wilkinson Prize for Numerical Software
- ▶ Widely used
 - ▶ Latest version 3.3.10
 - ▶ Includes threads, OpenMP, MPI versions, new architecture support
 - ▶ Layout constraints from users/apps + network differences are hard to support

More FFTW

- ▶ C library for real and complex FFTs (arbitrary size/dimensionality)
- ▶ Computational kernels (80% of code) automatically generated
- ▶ Self-optimizes for your hardware (picks best composition of steps) = portability + performance
- ▶ Exploits CPU-specific SIMD instructions (rewriting the code)
- ▶ FFTW implements many FFT algorithms: A planner picks the best composition by measuring the speed of different combination
- ▶ Free! Gnu!

FFTW selling points

- ▶ Speed. (Supports SSE/SSE2/Altivec)
- ▶ Both one-dimensional and multi-dimensional transforms.
- ▶ Arbitrary-size transforms. (Sizes with small prime factors are best, but FFTW uses $O(N \log N)$ algorithms even for prime sizes.)
- ▶ Fast transforms of purely real input or output data.
- ▶ Transforms of real even/odd data: the discrete cosine transform (DCT) and the discrete sine transform (DST).
- ▶ Efficient handling of multiple, strided transforms.
- ▶ Parallel transforms: parallelized code for platforms with SMP machines with some flavor of threads (e.g. POSIX) or OpenMP. Also MPI.
- ▶ Portable to any platform with a C compiler.
- ▶ Both C and Fortran interfaces.