

Outline

- Timeline
- Parallel graph algorithms
 - Applications
 - Designing parallel graph algorithms: DFS, BFS
 - Case studies:
 - A) Shortest Paths: Classic, Delta-stepping
 - B) Maximal Independent Sets: Luby's algorithm
 - C) Graph traversals: Breadth-first search
 - D) Strongly Connected Components

Graph reminder

- Define: Graph $G = (V, E)$: a set of vertices V and a set of edges E between vertices
- $n = |V|$ (number of vertices)
- $m = |E|$ (number of edges)
- $D = \text{diameter}$ (max #hops between any pair of vertices)
- Edges can be directed or undirected, weighted or not.
- They can even have attributes (i.e. semantic graphs)
- Sequences of edges $\langle u_1, u_2 \rangle, \langle u_2, u_3 \rangle, \dots, \langle u_{n-1}, u_n \rangle$ is a *path* from u_1 to u_n . Its *length* is the sum of its weights.

Many types of graphs

- Lines and trees
- Completely regular grids
- Planar graphs (no edges need cross)
- Low-dimensional Euclidean
- Power law graphs
- ...

Algorithms are not one-size-fits-all!

Applications (1)

- Routing in transportation networks (route planning with costs)
- Internet and the WWW
 - The world-wide web can be represented as a directed graph
 - Web search and crawl: traversal
 - Link analysis, ranking: Page rank
 - Document classification and clustering
 - Internet topologies (router networks) are naturally modeled as graphs

Applications (2): Large Graphs in Scientific Computing


- Graph partitioning: Dynamic load balancing in parallel simulations
- Problem size: as big as the sparse linear system to be solved or the simulation to be performed

Applications (3): Large-scale data analysis

- Graph abstractions are very useful to analyze complex data sets.
- Sources of data: simulations, experimental devices, the Internet, sensor networks
- Challenges: data size, heterogeneity, uncertainty, data quality
- Examples:
 - Astrophysics: massive datasets, temporal variations
 - Bioinformatics: data quality, heterogeneity
 - Social Informatics: new analytics challenges, data uncertainty



Applications (4): Large Graphs in Biology

- Graph Theoretical analysis of Brain Connectivity
- Whole genome assembly
 - 26 billion (8B of which are non-erroneous) unique k-mers (vertices) in the hexaploid wheat genome W7984 for k=51



Graph representations

Compressed sparse rows (CSR) = cache-efficient adjacency lists



Index into adjacency array	1	3	3	4	6
Adjacencies	1	3	2	2	4
Weights	12	26	19	14	7

(row pointers in CSR)

(column ids in CSR)

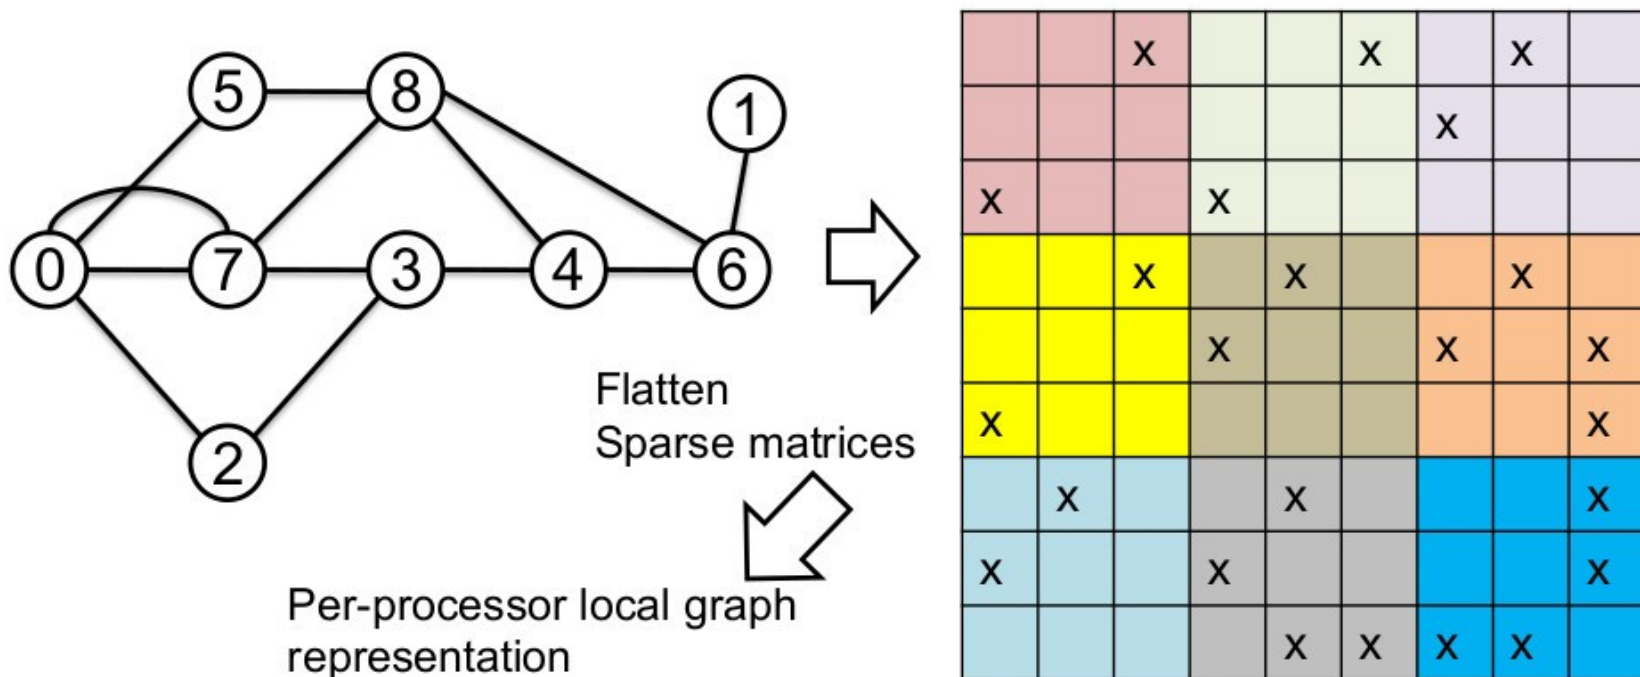
(numerical values in CSR)

Distributed graph representations

- Each processor stores the entire graph (“full replication”)
- Each processor stores n/p vertices and all adjacencies out of these vertices (“1D partitioning”)
 - How to create these “ p ” vertex partitions?
 - Graph partitioning algorithms: recursively optimize for conductance (edge cut/size of smaller partition)
 - Randomly shuffling the vertex identifiers ensures that edge count/processor are roughly the same

2D checkerboard distribution

- Consider a logical 2D processor grid ($p_r * p_c = p$) and the matrix representation of the graph
- Assign each processor a sub-matrix (i.e, the edges within the sub-matrix)



Memory Bandwidth Problems

- Many graph operations are
 - Computationally cheap (per node or edge)
 - Bad for locality
- Consider:
 - 323 million in US (fits in 32-bit int)
 - About 350 Facebook friends each
 - Compressed sparse row: about 450 GB
- What representation?

Memory Bandwidth Issues

- Adjacency Matrix
 - Pro: efficient for dense graphs
 - Con: wasteful for sparse case...
- Coordinates: Tuples: (i, j, w_{ij})
 - Pro: Easy to update
 - Con: Slow for multiply
- Adjacency list: Linked lists of adjacent nodes
 - Pro: Still easy to update
 - Con: May cost more to store than coordinates?

Distributed Graph Algorithm Design

- DFS and BFS components
- Distribute to processors
- Minimize communication

Graph traversal: Depth-first search (DFS)

```
procedure DFS(vertex v)
  v.visited = true
  previsit (v)
  for all v s.t.  $(v, w) \in E$ 
    if (!w.visited) DFS(w)
  postvisit (v)
```

Graph traversal:

Parallel Depth-first search (DFS)

- Each processor maintains a frontier of vertices
- As in sequential DFS, a frontier stores the subset of visited vertices, whose outgoing edges have not yet been explored.
- When a processor discovers a new vertex, it attempts to visit the vertex by using an *atomic read-modify-write* operation and, if it succeeds, adds the vertex to its frontier.
- Must perform load balancing to keep all the processors busy.
- A naive approach to this end would be to generate one thread for each vertex in the frontier.

Graph traversal:

Serial Breadth-first search (BFS)

Push seed node onto queue and mark

While queue nonempty:

- Pop node from queue


- Visit node

- Push unmarked neighbors on queue

- Mark all neighbors


Graph traversal: Parallel Breadth-first search (BFS)

- Expand current frontier (level-synchronous approach, suited for low diameter graphs)
 - $O(D)$ parallel steps
 - Adjacencies of all vertices in current frontier are visited in parallel



Graph traversal: Parallel Breadth-first search (BFS)

- Stitch together multiple concurrent traversals (Ullman-Yannakakis approach, suited for high-diameter graphs)
 - path-limited searches from “super vertices”
 - All-pairs shortest path between “super vertices”



Serial BFS: Bottom up

- Classical (top-down) algorithm is optimal in worst case, but pessimistic for low-diameter graphs
- Direction Optimization:
 - Switch from top-down to bottom-up search
 - When the majority of the vertices are discovered.
- Top down: For all v in frontier attempt to “parent” all neighbors(v)
- Bottom up: For all v in unvisited find any parent (neighbor(v) in frontier)

Serial BFS: Bottom up

Set $d[v] = \infty$ for all vertices

Set $d[s] = 0$ for seed s

Until d stops changing

For each $u \in V$

$$d[u] = \min(d[u], \min_{w \in N(u)} d[w] + 1)$$

Graph Algorithm Case Studies

- Shortest Paths: Delta-stepping, Floyd-Warshall
- Maximal Independent Sets: Luby's algorithm
- Graph traversals: Breadth-first search
- Strongly Connected Components

Parallel Single-source Shortest Paths (SSSP) algorithms

- Famous serial algorithms:
 - Bellman-Ford : label correcting - works on any graph
 - Dijkstra : label setting – requires nonnegative edge weights
- No known PRAM algorithm that runs in sub-linear time and $O(m+n \log n)$ work
- Ullman-Yannakakis randomized approach
- Meyer and Sanders, Δ - stepping algorithm

Classic Dijkstra

- Dequeue *closest point* to frontier, expand frontier
- Update priority queue of distances (in parallel)
- Repeat

Classic Bellman-Ford

- Initialize $d[u]$ with distance over-estimates to source

$$d[s]=0$$

Repeatedly relax $d[u] := \min_{(v,u) \in E} d[v] + w(v,u)$

- Converges (eventually) as long as all nodes visited repeatedly, updates are atomic

Δ - stepping algorithm

- Label-correcting algorithm: Can relax edges from unsettled vertices also
- “approximate bucket implementation of Dijkstra”
- For random edge weights $[0,1]$, runs in $O(n + m + D \cdot L)$ where $L = \max$ distance from source to any node
- Vertices are ordered using buckets of width Δ
- Each bucket may be processed in parallel
- Basic operation: Relax ($e(u,v)$):
 - $d(v) = \min \{ d(v), d(u) + w(u, v) \}$
- $\Delta < \min w(e)$: Degenerates into Dijkstra
- $\Delta > \max w(e)$: Degenerates into Bellman-Ford

Maximal Independent Sets (MIS)

- Graph with vertices $V = \{1, 2, \dots, n\}$
- A set S of vertices is *independent* if no two vertices in S are neighbors.
- An independent set S is *maximal* if it is impossible to add another vertex and stay independent
- An independent set S is *maximum* if no other independent set has more vertices
- Finding a maximum independent set is intractably difficult (NP-hard)
- Finding a maximal independent set is easy, at least on one processor.

Sequential (Greedy) Maximal Independent Set Algorithm

S = empty set;

for vertex $v = 1$ to n

 if (v has no neighbor in S)

 add v to S

Luby's Parallel Algorithm

$S = \text{empty set}; C = V;$

while C is not empty

 label each v in C with a random $r(v)$;

 for all v in C *in parallel*

 if $r(v) < \min(r(\text{neighbors of } v))$

 move v from C to S ;

 remove neighbors of v from C

- “Probably” finishes in $O(\log n)$ rounds

Strongly Connected Components

- Strongly connected components (SCCs): all maximal strongly connected sub-graphs in a large directed graph.
- Sequential algorithm: use depth-first search (Tarjan); work= $O(m+n)$ for $m=|E|$, $n=|V|$, but DFS seems to be inherently sequential.
- Parallel algorithm: divide-and-conquer and BFS (Fleischer et al.); worst-case span $O(n)$ but good in practice on many graphs.

Fleischer/Hendrickson/Pinar algorithm

- Partition the given graph into three disjoint subgraphs
- Each can be processed independently/recursively
- $FW(v)$: vertices reachable from vertex v (ForWard).
- $BW(v)$: vertices from which v is reachable (BackWard).

Fleischer/Hendrickson/Pinar algorithm

- 1) Select maximally connected node, called the *pivot*
- 2) Find all vertices that can be reached from the pivot (descendant (D))
- 3) Find all vertices that can reach the pivot (predecessor (P))
- 4) Intersection of those two sets is an SCC ($S = P \cap D$)
- 5) Now have three distinct sets leftover ($D \setminus S$), ($P \setminus S$), and remainder (R)

$V \setminus X$ denotes the subset of vertices in V which are not in a subset X

Fleischer/Hendrickson/Pinar algorithm

procedure FW-BW(V)

if $V = \emptyset$ then return \emptyset

Select a pivot $u \in V$

$D \leftarrow \text{BFS}(G(V, E(V)), u)$

$P \leftarrow \text{BFS}(G(V, E'(V)), u)$

$R \leftarrow (V \setminus (P \cup D))$

$S \leftarrow (P \cap D)$

new task do FW-BW($D \setminus S$)

new task do FW-BW($P \setminus S$)

new task do FW-BW(R)