

Midterm Review

Flynn's Taxonomy

*Mike Flynn, "Very High-Speed Computing Systems,"
Proc. of IEEE, 1966*

- **Single Instruction Single Data**
- **Single Instruction Multiple Data**
 - Array Processor
 - Vector Processor
 - GPU
- **Multiple Instruction Single Data**
 - Closest form: systolic array processor, streaming processor
- **Multiple Instruction Multiple Data**
 - Multiprocessor
 - Multithreaded processor

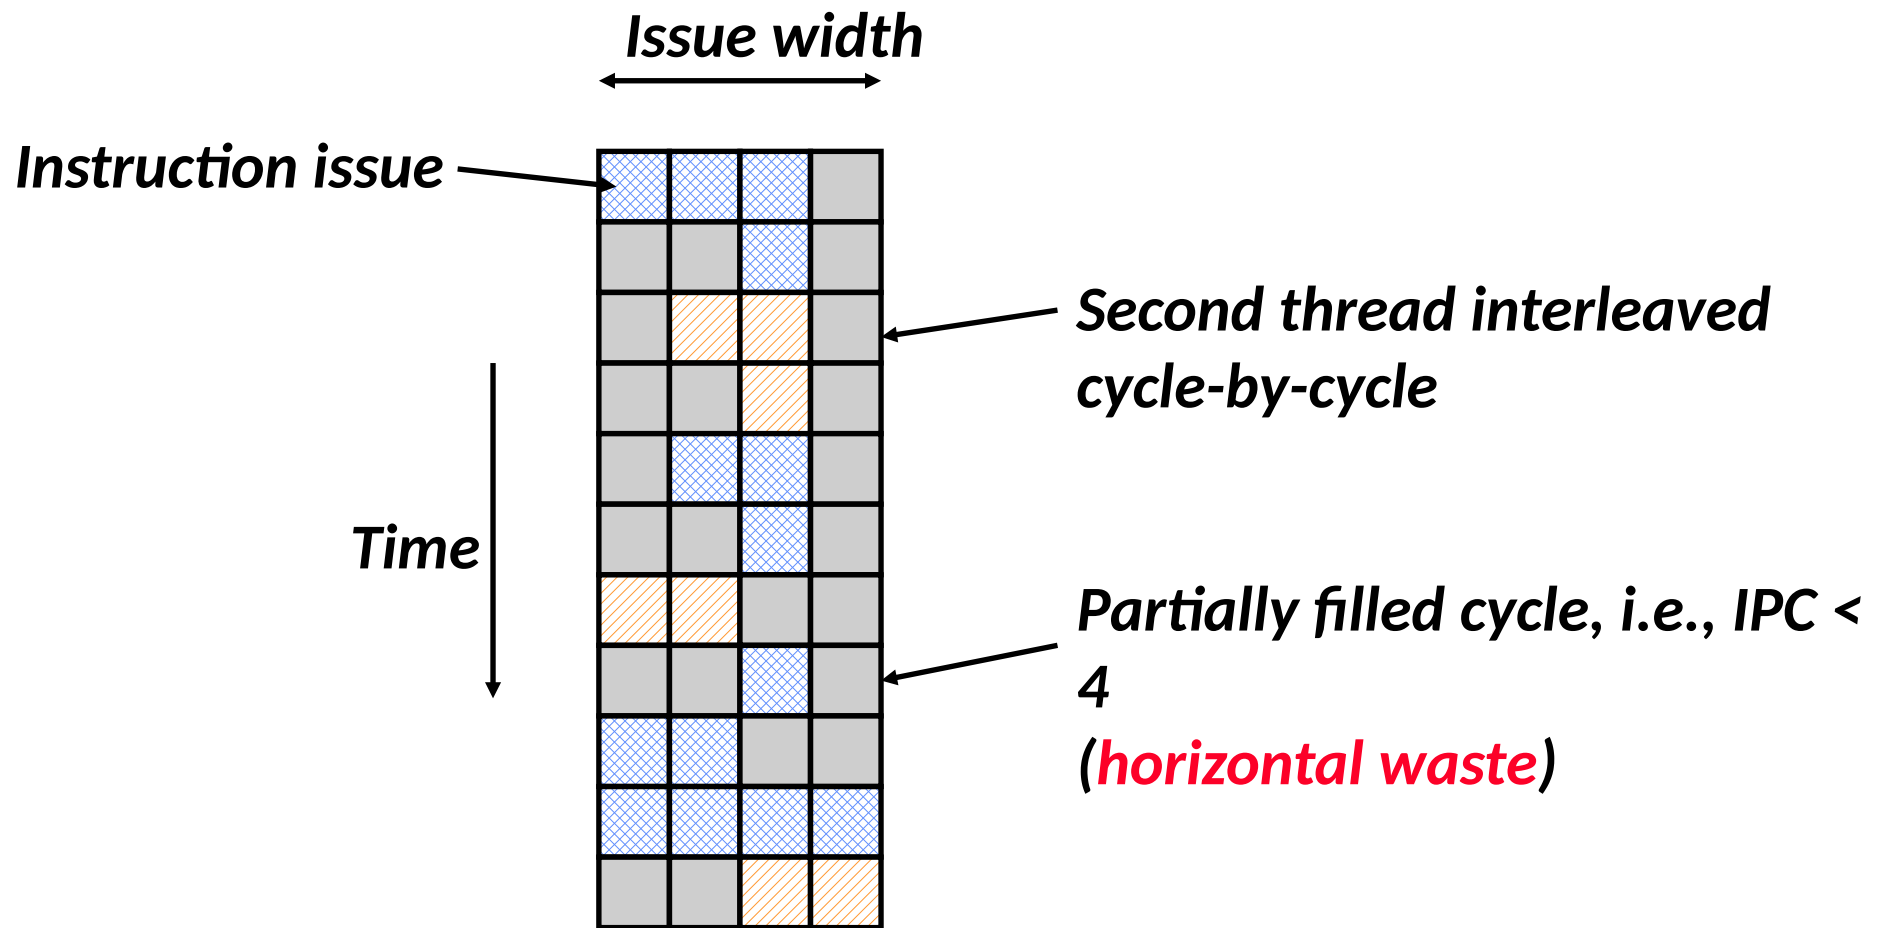
Instruction Level Parallelism (ILP)

- Improve processor performance by having multiple processor components or functional units simultaneously executing instructions.
- Pipelining - functional units are arranged in *stages*.
- Multiple issue - multiple instructions *might* be simultaneously initiated with resource duplication.
 - VLIW – Very long instruction word - functional units are scheduled at compile time (static scheduling).
 - Superscalar - functional units are scheduled at run-time (dynamic scheduling)

Multithreading

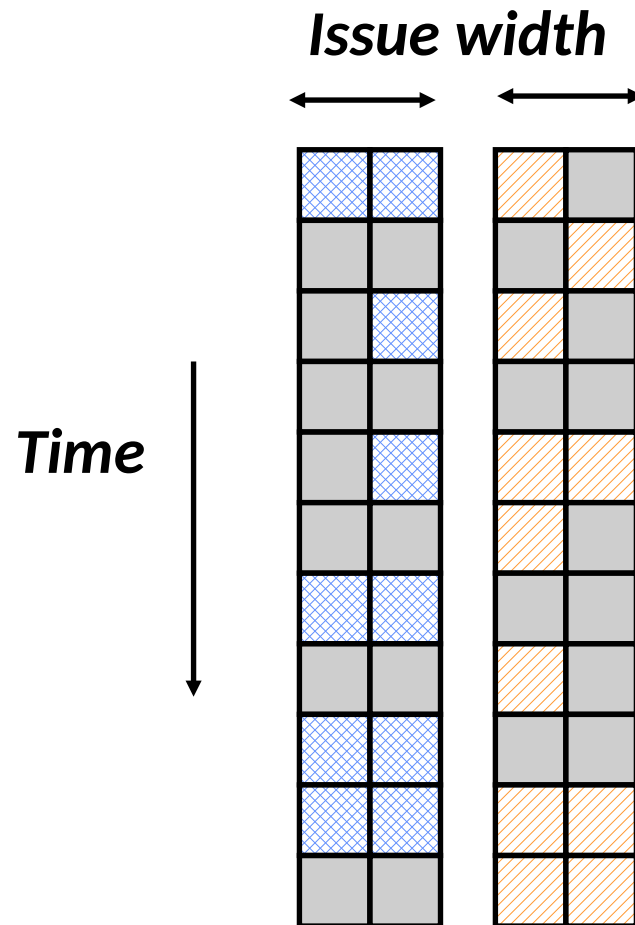
- Difficult to continue to extract instruction-level parallelism (ILP) from a single sequential thread of control
- Many workloads can make use of thread-level parallelism (TLP)
 - TLP from multiprogramming (run independent sequential jobs)
 - TLP from multithreaded applications (run one job faster using parallel threads)
- Multithreading uses TLP to improve utilization of a single processor

Vertical Multithreading



- Cycle-by-cycle interleaving removes vertical waste, but leaves some horizontal waste

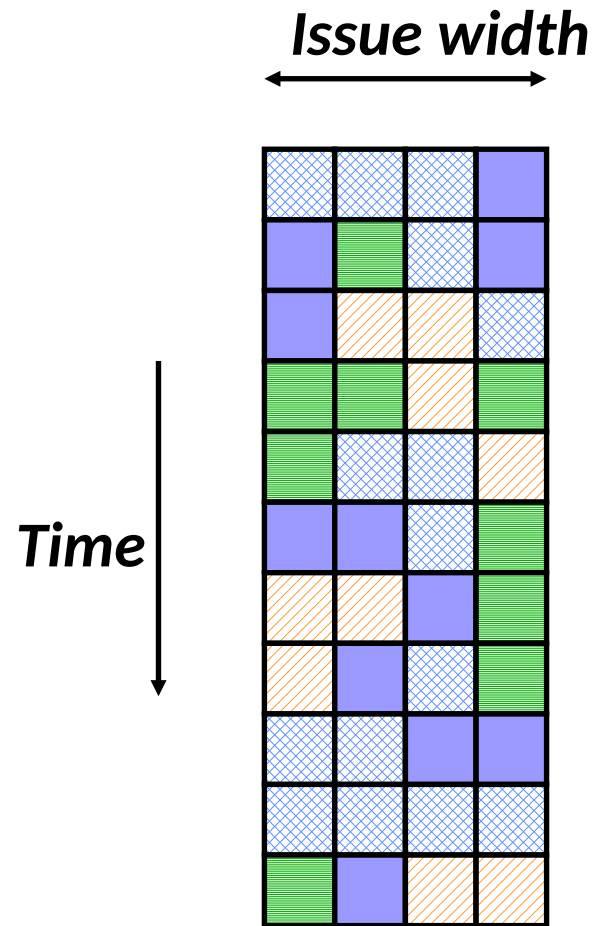
Chip Multiprocessing (CMP)



- What is the effect of splitting into multiple processors?
 - reduces horizontal waste,
 - leaves some vertical waste, and
 - puts upper limit on peak throughput of each thread.

Ideal Superscalar Multithreading

[Tullsen, Eggers, Levy, UW, 1995]



- Interleave multiple threads to multiple issue slots with no restrictions

Importance of Data Parallelism

- GPUs are designed for graphics
 - Highly parallel tasks
- Process *independent* vertices & fragments
 - No shared or static data
 - No read-modify-write buffers
- Data-parallel processing
 - GPU architecture is ALU-heavy
 - Performance depends on *arithmetic intensity*
 - Computation / Bandwidth ratio
 - Hide memory latency with more computation

Arithmetic Intensity

Lots of operations per word transferred

Graphics pipeline

Vertex

Bandwidth: 1 triangle = 32 bytes;

Operations: 100-500 f32-ops / triangle

Rasterization

Create 16-32 fragments per triangle

Fragment

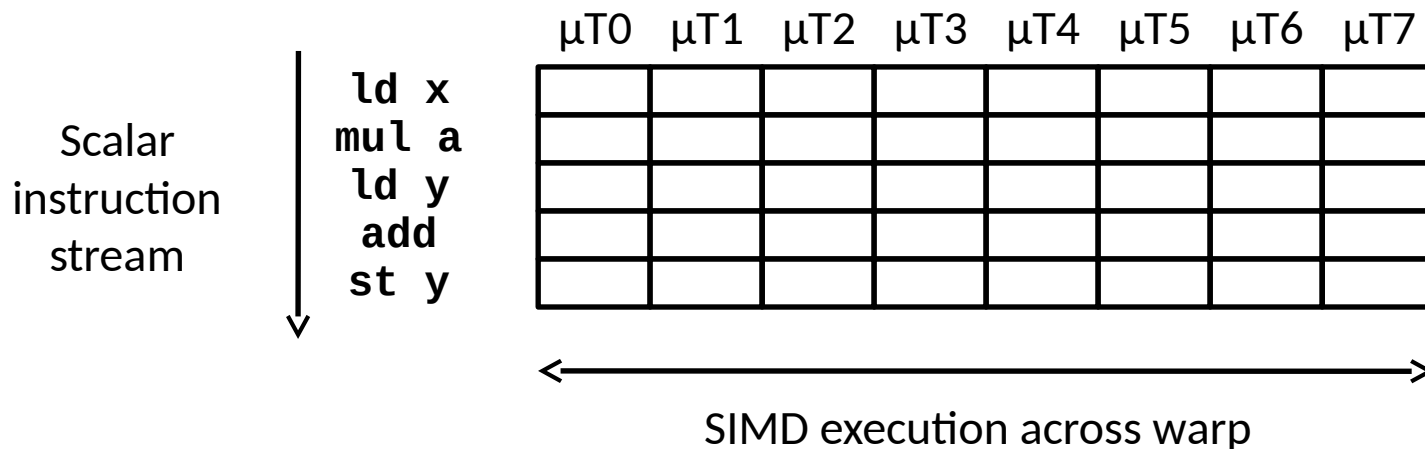
Bandwidth: 1 fragment = 10 bytes

Operations: 300-1000 i8-ops/fragment

Courtesy of Pat Hanrahan

“Single Instruction, Multiple Thread”

- GPUs use a **SIMT** model, where individual scalar instruction streams for each CUDA thread are grouped together for SIMD execution on hardware (Nvidia groups 32 CUDA threads into a *warp*)

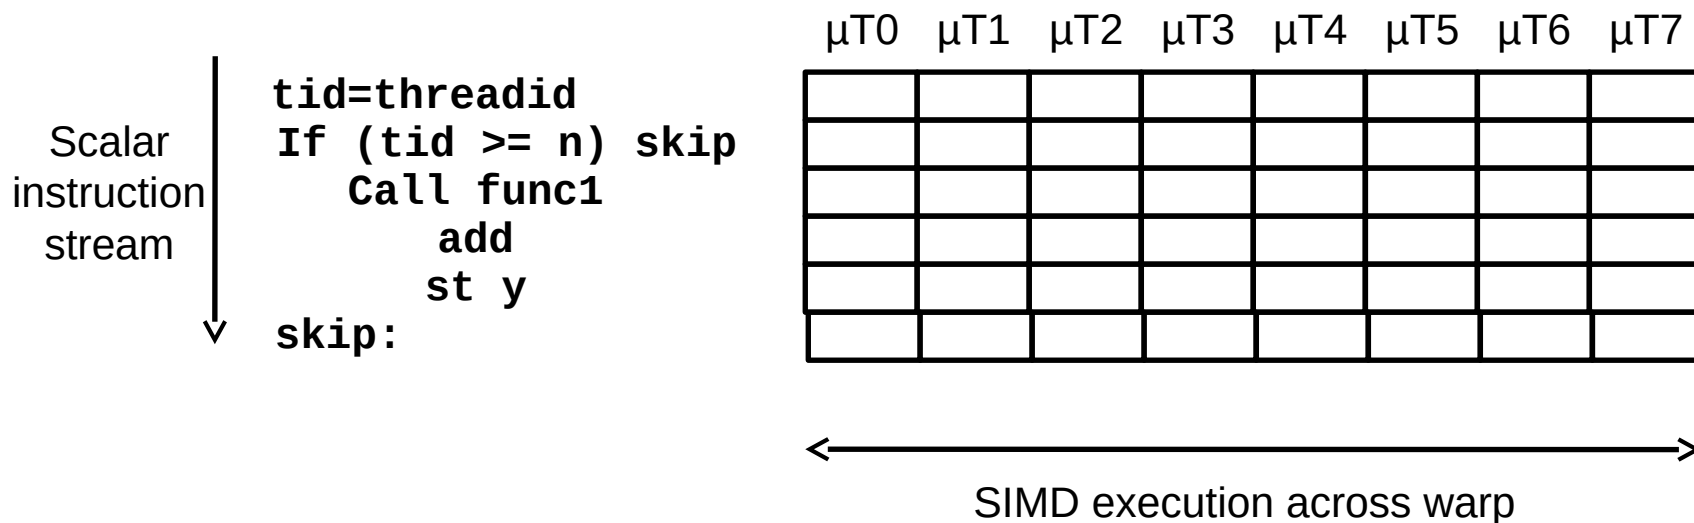


Implications of SIMT model

- All “vector” loads and stores are scatter-gather, as individual μ threads perform scalar loads and stores
 - GPU adds hardware to dynamically coalesce individual μ thread loads and stores to mimic vector loads and stores
- Every μ thread has to perform stripmining calculations redundantly (“am I active?”) as there is no scalar processor equivalent
- Illusion of many independent threads
- But for efficiency, programmer must try and keep μ threads aligned in a SIMD fashion
 - Try and do unit-stride loads and store so memory coalescing kicks in
 - Avoid branch divergence so most instruction slots execute useful work and are not masked off

Conditionals in SIMT model

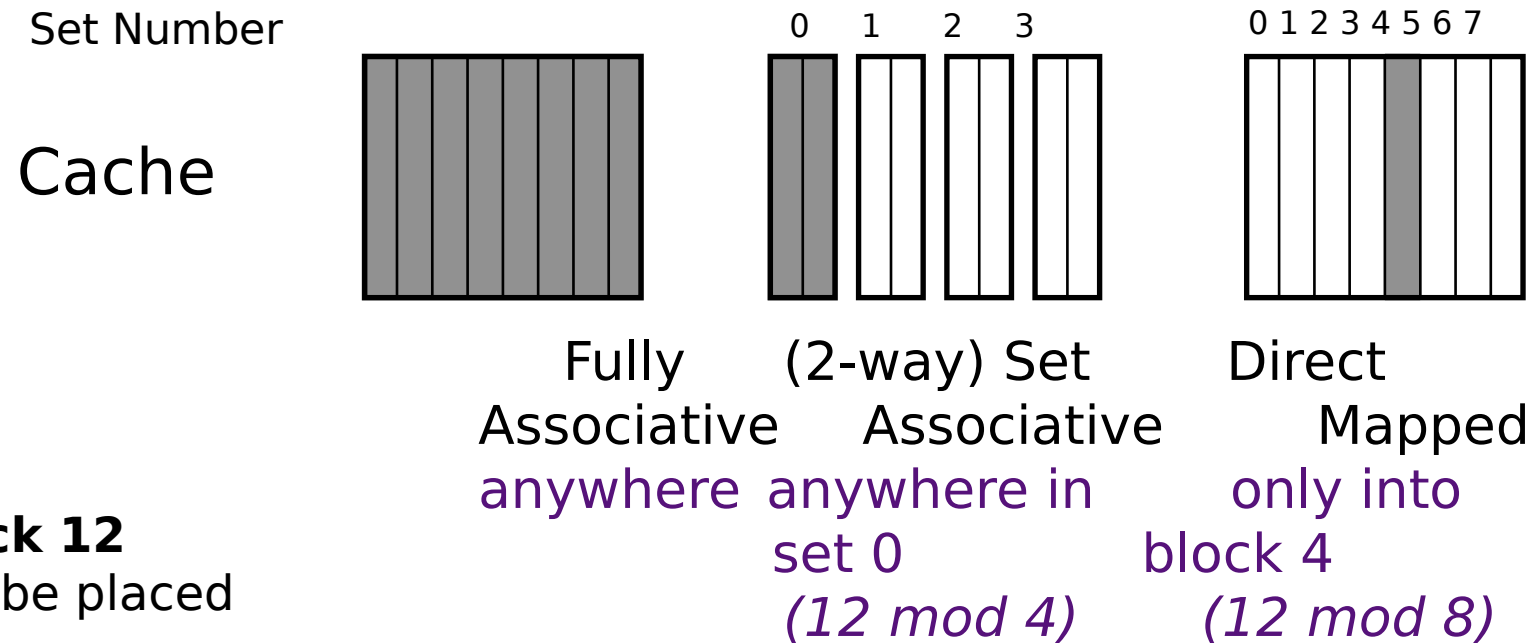
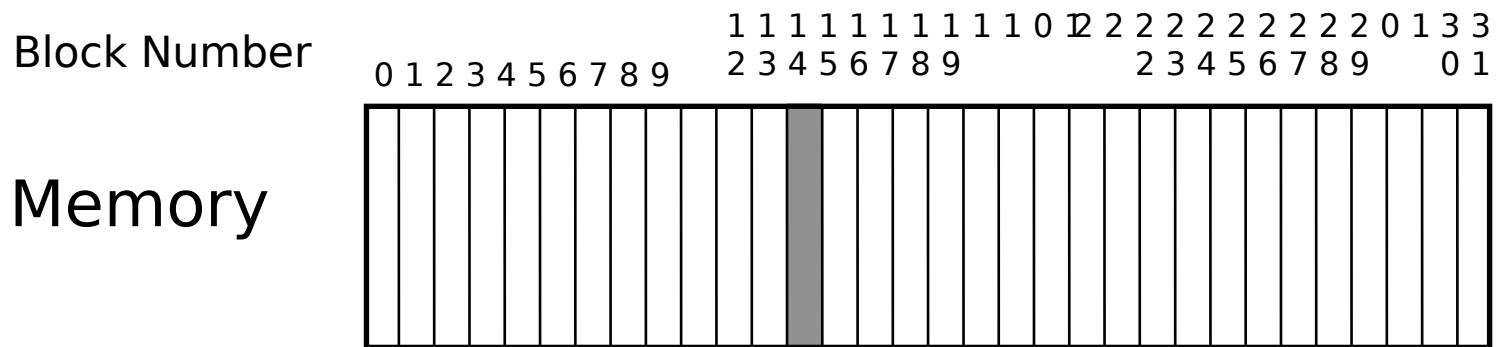
- Simple if-then-else are compiled into *predicated execution*, equivalent to vector masking
- More complex control flow compiled into branches
- How to execute a vector of branches?



Cache Organizations

- “Fully Associative”: Block can go anywhere
 - Note: No Index field, but one comparator/block
- “Direct Mapped”: Block goes one place
 - Note: Only 1 comparator
 - Number of sets = number blocks
- “N-way Set Associative”: N places for a block
 - Number of sets = number of blocks / N
 - N comparators

Placement Policy



block 12
can be placed

Replacement Policy

In an associative cache, which block from a set should be evicted when the set becomes full?

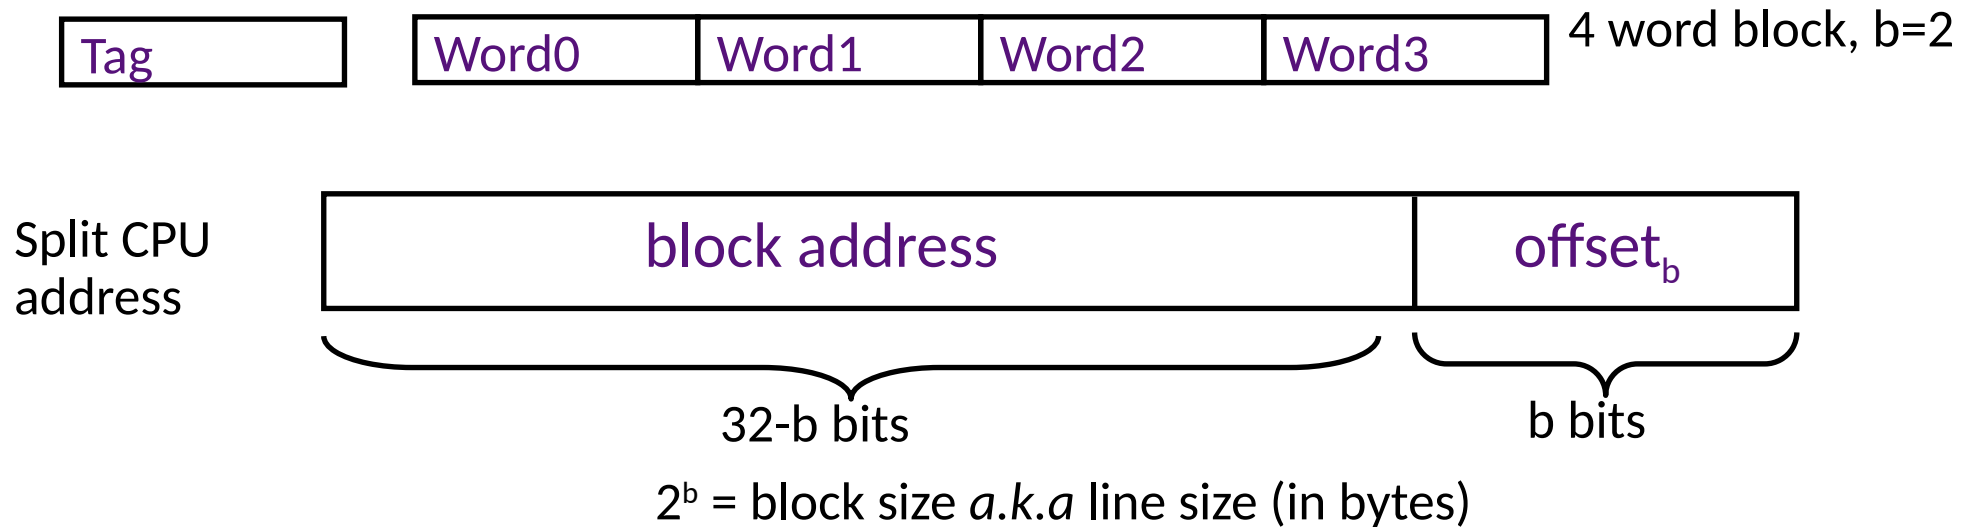
- Random
- Least-Recently Used (LRU)
 - LRU cache state must be updated on every access
 - true implementation only feasible for small sets (2-way)
 - pseudo-LRU binary tree often used for 4-8 way
- First-In, First-Out (FIFO) a.k.a. Round-Robin
 - used in highly associative caches
- Not-Most-Recently Used (NMRU)
 - FIFO with exception for most-recently used block or blocks

This is a second-order effect. Why?

Replacement only happens on misses

Block Size and Spatial Locality

Block is unit of transfer between the cache and memory



Larger block size has distinct hardware advantages

- less tag overhead
- exploit fast burst transfers from DRAM
- exploit fast burst transfers over wide busses

What are the disadvantages of increasing block size (assuming constant memory size)?

Fewer blocks => more conflicts. Can waste bandwidth.

False Sharing

state	line addr	data0	data1	...	dataN
-------	-----------	-------	-------	-----	-------

A cache line contains more than one word

Cache-coherence is done at the line-level and *not* word-level

Suppose M_1 writes word_i and M_2 writes word_k and both words have the same line address.

What can happen?

Coherency misses

- 1) **True sharing misses** arise from the communication of data through the cache coherence mechanism
 - Invalidates due to 1st write to shared line
 - Reads by another CPU of modified line in different cache
 - Miss would still occur if line size were 1 word, i.e., truly shared
- 2) **False sharing misses** when a line is invalidated because a word in the line, other than the one being read, is written into
 - Invalidation does not cause a new value to be communicated, but only causes an extra cache miss
 - Line is shared, but no word in line is *actually shared*
 - miss would not occur if line size were 1 word

Blocked (Tiled) Matrix Multiply

```
// implements C = C + A*B
// A,B,C to be N-by-N matrices of b-by-b subblocks
for i = 1 to n
  for j = 1 to n
    for k = 1 to n
      // Matrix multiply over blocks of size b = n / N
      C(i,j) = C(i,j) + A(i,k) * B(k,j)
```

if $n \% b \neq 0$ you need to have code for the “edges/fringes”

$2n^2$ to read and write each block of C once: $2N^2 * b^2 = 2n^2$

$N * n^2$ to read each block of A N^3 times: $N^3 * b^2 = N^3 * (n/N)^2$

$N * n^2$ to read each block of B N^3 times: $N^3 * b^2 = N^3 * (n/N)^2$

Computational Intensity, $CI = f / m = 2n^3 / ((2N + 2) * n^2) \approx n / N = b$ for large n

Take-Aways

- Matrix vector and matrix matrix multiplication key (linear algebra)
- Matrix vector multiplication
 - Opportunities for better parallelism and use special instructions, Fused Multiply Add, etc.
 - Limited by bandwidth ($O(2n^2)$ flops on $O(n^2)$ data)
- Matrix matrix multiplication
 - Can improve computational intensity $O(2n^3)$ flops on $O(3n^2)$ data
 - Tiling (aka blocking)
- Optimizations in practice
 - Expose parallelism to compiler (SIMD, ILP, prefetching)
 - Help compiler with register use
 - Help hardware (data layout) with cache locality (temporal and spatial)
- Optimized libraries (BLAS) exist

Optimizing in Practice

- Tiling for registers
 - loop unrolling, use of named “register” variables
- Tiling for multiple levels of cache and TLB
- Exploiting fine-grained parallelism in processor
 - superscalar; pipelining
- Complicated compiler interactions (flags)
- Hard to do by hand

Note on Matrix Storage

- A matrix is a 2-D array of elements, but memory addresses are “1-D”
- Conventions for matrix layout
 - by column, or “column major” (Fortran default); $A(i,j)$ at $A+i+j*n$
 - by row, or “row major” (C default) $A(i,j)$ at $A+i*n+j$
 - recursive (later)
- Intense interaction with cache line size

Tips on Tuning

“We should forget bout small efficiencies, say about 97% of the time: premature optimization is the root of all evil.”– C.A.R. Hoare (quoted by Donald Knuth)

- Tradeoff: speed vs readability, debuggability, maintainability...
- Only optimize when needful
- Go for low-hanging fruit first: data layouts, libraries, compiler flags
- Concentrate on the bottleneck
- Concentrate on inner loops
- Get correctness (and a test framework) first

Loop Unrolling

- Expose instruction-level parallelism and reduce control overhead

```
float f0 = filter[0], f1 = filter[1], f2 = filter[2];
```

```
float s0 = signal[0], s1 = signal[1], s2 = signal[2];
```

```
*res++ = f0*s0 + f1*s1 + f2*s2;
```

```
do {
```

```
    signal += 3;
```

```
    s0 = signal[0];
```

```
    res[0] = f0*s1 + f1*s2 + f2*s0;
```

```
    s1 = signal[1];
```

```
    res[1] = f0*s2 + f1*s0 + f2*s1;
```

```
    s2 = signal[2];
```

```
    res[2] = f0*s0 + f1*s1 + f2*s2;
```

```
    res += 3;
```

```
} while( ... );
```


Basic Types of Synchronization: Mutexes

Mutexes -- mutual exclusion aka locks

threads are working mostly independently

need to access common data structure

```
lock *l = alloc_and_init();  /* shared */  
acquire(l);  
    access data  
release(l);
```

Locks only affect processors using them:

If a thread accesses the data without doing the acquire/release, locks by others will not help

Java, C++, and other languages have lexically scoped synchronization, i.e., synchronized methods/blocks (In Java: synchronized (l) {...})

Can't forgot to say "release"

Semaphores (a signaling mechanism) generalize locks to allow k threads simultaneous access; good for limited resources.

Unlike in a mutex, a semaphore can be decremented by another process (a mutex can only be unlocked by its owner)

Example Problem: Arithmetic Intensity

```
for (int i=1; i<n; i++)  
    for (int j=1; j<n; j++) {  
        u(i,j) = u(i,j) - 0.125*( u(i+1,j) + u(i-1,j) +  
            u(i+1,j-1) + u(i,j-1) + u(i-1,j-1) +  
            u(i+1,j+1) + u(i-1,j+1) + u(i,j+1));  
    };
```

What is OpenMP?

- OpenMP = Open specification for Multi-Processing
 - openmp.org – Talks, examples, forums, etc.
 - Specification controlled by the Architecture Review Board (ARB)
- Motivation: capture common usage and simplify programming
- OpenMP Architecture Review Board (ARB)
 - A nonprofit organization that controls the OpenMP Spec
 - Latest spec: OpenMP 5.0 (Nov. 2018)
- High-level API for programming in C/C++ and Fortran
 - Preprocessor (compiler) directives (~ 80%)
`#pragma omp construct [clause [clause ...]]`
 - Library Calls (~ 19%)
`#include <omp.h>`
 - Environment Variables (~ 1%)
all caps

The OpenMP Common Core:

Most OpenMP programs use these 9+10 items

<code>#pragma omp parallel</code>	Parallel region, teams of threads, structured block, interleaved execution across threads
<code>setenv OMP_NUM_THREADS N</code>	Internal control variables. Setting the default number of threads with an environment variable
<code>#pragma omp barrier</code> <code>#pragma omp critical</code>	Synchronization and race conditions. Revisit interleaved execution.
<code>#pragma omp for</code> <code>#pragma omp parallel for</code>	Worksharing, parallel loops, loop carried dependencies
<code>#pragma omp single</code>	Workshare with a single thread
<code>#pragma omp task</code> <code>#pragma omp taskwait</code>	Tasks including the data environment for tasks.

The OpenMP Common Core:

Most OpenMP programs use these 9+10 items

`int omp_get_thread_num()`
`int omp_get_num_threads()`

Create threads with a parallel region and split up the work using the number of threads and thread ID

`double omp_get_wtime()`

Speedup: False Sharing and other performance issues

`reduction(op:list)`

Reductions of values across a team of threads

`schedule(dynamic [,chunk])`
`schedule (static [,chunk])`

Loop schedules, loop overheads and load balance

`private(list)`, `firstprivate(list)`, `shared(list)`

Data environment

`nowait`

Disabling implied barriers on workshare constructs, the high cost of barriers, and the flush concept (but not the flush directive)

OpenMP basic syntax

- Most of the constructs in OpenMP are compiler directives.
- Most OpenMP* constructs apply to a “structured block”.
 - Structured block: a block of one or more statements with one point of entry at the top and one point of exit at the bottom.
 - It’s OK to have an exit() within the structured block.

	C and C++	Fortran
Compiler directive	<code>#pragma omp construct [clause [clause] ...]</code>	<code>!\$OMP construct [clause [clause] ...]</code>
Example	<code>#pragma omp parallel private(x) { ... }</code>	<code>!\$OMP PARALLEL ... !\$OMP END PARALLEL</code>
Function prototype	<code>#include <omp.h></code>	<code>use OMP_LIB</code>

Thread creation: Parallel regions

- You create threads in OpenMP with the **parallel** construct.
- Example, To create a 4 thread Parallel region:

```
double A[1000];  
omp_set_num_threads(4);  
#pragma omp parallel  
{  
    int ID = omp_get_thread_num();  
    pahrump(ID,A);  
}
```

- Each thread calls pahrump(ID,A) for ID = 0 to 3
- OpenMP does not have to delete the threads here; it can create a thread pool and reuse it (implementation dependent)
- This is in contrast to pthreads, which is *imperative* (the library will kill all threads by definition during pthread_join)

Example problem

```
#define N 1000
Extern struct data member [ N ] ;
Extern int isgood (int i ) ;
int good_members [ N ] ;
int pos = 0 ;

void find_good_members ( )
{
#pragma omp parallel for
    For (i=0; i< N; i++) {
        If ( is_good ( i ) ) {
            good_members [ pos ] = i ;
#pragma omp atomic
            pos++
        }
    }
}
```


Why such poor scaling?

False sharing

- If independent data elements happen to sit on the same cache line, each update will cause the cache lines to “slosh back and forth” between threads ... This is called “*false sharing*”.
- If you promote scalars to an array to support creation of an SPMD program, the array elements are contiguous in memory and hence share cache lines ... Results in poor scalability.
- Solution: Pad arrays so elements you use are on distinct cache lines.

Snoopy bus protocol (reprise)

- Basic idea:
 - Broadcast operations on memory bus
 - Cache controllers “snoop” on all bus transactions
 - Memory writes induce serial order
 - Act to enforce coherence (invalidate, update, etc)
- Problems:
 - Bus bandwidth limits scaling
 - Contending writes are slow
- There are other protocol options (e.g. directory-based).
 - But usually give up on *full* sequential consistency.

Synchronization: critical section

```
float res;
```

```
#pragma omp parallel
```

```
{    float B;  int i, id, nthrds;
```

```
    id = omp_get_thread_num();
```

```
    nthrds = omp_get_num_threads();
```

```
    for(i=id;i<niters;i+=nthrds){
```

```
        B = big_job(i);
```

```
#pragma omp critical
```

```
    res += consume (B); // Thread wait: One at a time...
```

```
}
```

```
}
```

Synchronization: Barrier

- Barrier: a point in a program all threads must reach before any threads are allowed to proceed.
- It is a “stand alone” pragma meaning it is not associated with user code ... it is an executable statement.

```
double Arr[8], Brr[8];
int numthrds;
omp_set_num_threads(8)
#pragma omp parallel
{   int id, nthrds;
    id = omp_get_thread_num();
    nthrds = omp_get_num_threads();
    if (id==0) numthrds = nthrds;
    Arr[id] = big_ugly_calc(id, nthrds);
#pragma omp barrier
    Brr[id] = really_big_and_ugly(id, nthrds, Arr);
}
```

Reduction

- OpenMP reduction clause: **reduction** (op : list)
 - Inside a parallel construct:
 - A local copy of each list variable is made and initialized depending on the “op” (e.g. 0 for “+”)
 - Updates occur on the *local copy*
 - Local copies are reduced into a single value and combined with the original global value.
- The variables in “list” must be shared in the enclosing parallel region.

```
double ave=0.0, A[MAX];  int i;
```

```
#pragma omp parallel for reduction (+:ave)
```

```
for (i=0;i< MAX; i++) {
```

```
    ave + = A[i];
```

```
} // implicit barrier here
```

```
ave = ave/MAX;
```

Takeaway

- Programming shared memory machines
 - May allocate data in large shared region without too many worries about where
 - Memory hierarchy is critical to performance
 - Even more so than on uniprocessors, due to coherence traffic
 - For performance tuning, watch sharing (both true and false)
- Semantics
 - Need to lock access to shared variable for read-modify-write
 - Sequential consistency is the natural semantics
 - Write race-free programs to get this
 - Architects worked hard to make this work
 - Caches are coherent with buses or directories
 - No caching of remote data on shared address space machines
 - But compiler and processor may still get in the way
 - Non-blocking writes, read prefetching, code motion...
 - Avoid races or use machine-specific fences carefully

Network Design

- **Topology** (how things are connected)
 - Crossbar; ring; 2-D, 3-D, higher-D mesh or torus; hypercube; tree; butterfly; perfect shuffle, dragon fly, ...
- **Routing algorithm:**
 - Example in 2D torus: all east-west then all north-south (avoids deadlock).
- **Switching strategy:**
 - Circuit switching: full path reserved for entire message, like the telephone.
 - Packet switching: message broken into separately-routed packets, like the post office, or internet
- **Flow control** (what if there is congestion):
 - Stall, store data temporarily in buffers, re-route data to other nodes, tell source node to temporarily halt, discard, etc.

Performance Properties of a Network

- **Diameter**: the maximum (over all pairs of nodes) of the shortest path between a given pair of nodes.
- **Latency**: delay between send and receive times
 - Latency tends to vary widely across architectures
 - Vendors often report **hardware latencies** (wire time)
 - Application programmers care about **software latencies** (user program to user program)
- **Observations**:
 - Latencies differ by 1-2 orders across network designs
 - Software/hardware overhead at source/destination dominate cost (1s-10s usecs)
 - Hardware latency varies with distance (10s-100s nsec per hop) but is small compared to overheads
- Latency is key for programs with many small messages

Latency and Bandwidth Model

- Time to send message of length n is roughly

$$\begin{aligned}\text{Time} &= \text{latency} + n * \text{cost_per_word} \\ &= \text{latency} + n / \text{bandwidth}\end{aligned}$$

- Topology is assumed irrelevant
- Often called “ α - β model” and written

$$\text{Time} = \alpha + n * \beta$$

- Usually $\alpha \gg \beta \gg \text{time per flop}$.
 - One long message is cheaper than many short ones.

$$\alpha + n * \beta \ll n * (\alpha + 1 * \beta)$$

- Can do hundreds or thousands of flops for cost of one message.
- Lesson: *Need large computation-to-communication ratio to be efficient*

Programming Distributed Memory Machines with Message Passing

- Overview of MPI
- Basic send/receive use
- Non-blocking communication
- Collectives

Message Passing Libraries

- All communication, synchronization require subroutine calls
 - *No shared variables*
 - Program run on a single processor just like any uniprocessor program, except for calls to message passing library
- Subroutines for
 - Communication
 - Pairwise or point-to-point: Send and Receive
 - Collectives all processor get together to
 - Move data: Broadcast, Scatter/gather
 - Compute and move: sum, product, max, prefix sum, ... of data on many processors
 - Synchronization
 - Barrier
 - No locks because there are no shared variables to protect
 - Inquiries
 - How many processes? Which one am I? Any messages waiting?

MPI Basic Concepts

- Processes can be collected into groups
- Each message is sent in a context, and must be received in the same context
 - Provides necessary support for libraries
- A group and context together form a communicator
- A process is identified by its rank in the group associated with a communicator
- There is a default communicator whose group contains all initial processes, called **MPI_COMM_WORLD**

Blocking and Non-blocking Communication

- So far we have been using *blocking* communication:
 - MPI_Recv does not complete until the buffer is full (available for use).
 - MPI_Send does not complete until the buffer is empty (available for use).
- Completion depends on size of message and amount of system buffering.

What MPI Functions are in Use?

For simple applications, these are common:

- Point-to-point communication
 - MPI_Irecv, MPI_Isend, MPI_Wait, MPI_Send, MPI_Recv
- Startup, Shutdown
 - MPI_Init, MPI_Finalize
- Information on the processes
 - MPI_Comm_rank, MPI_Comm_size, MPI_Get_processor_name
- Collective communication
 - MPI_Allreduce, MPI_Bcast, MPI_Allgather

MPI Collective Routines

- Many Routines: **Allgather, Allgatherv, Allreduce, Alltoall, Alltoallv, Bcast, Gather, Gatherv, Reduce, Reduce_scatter, Scan, Scatter, Scatterv**
- **All** versions deliver results to all participating processes.
- **V** versions allow the hunks to have variable sizes.
- **Allreduce, Reduce, Reduce_scatter, and Scan** take both built-in and user-defined combiner functions.
- MPI-2 adds **Alltoallw, Exscan**, intercommunicator versions of most routines

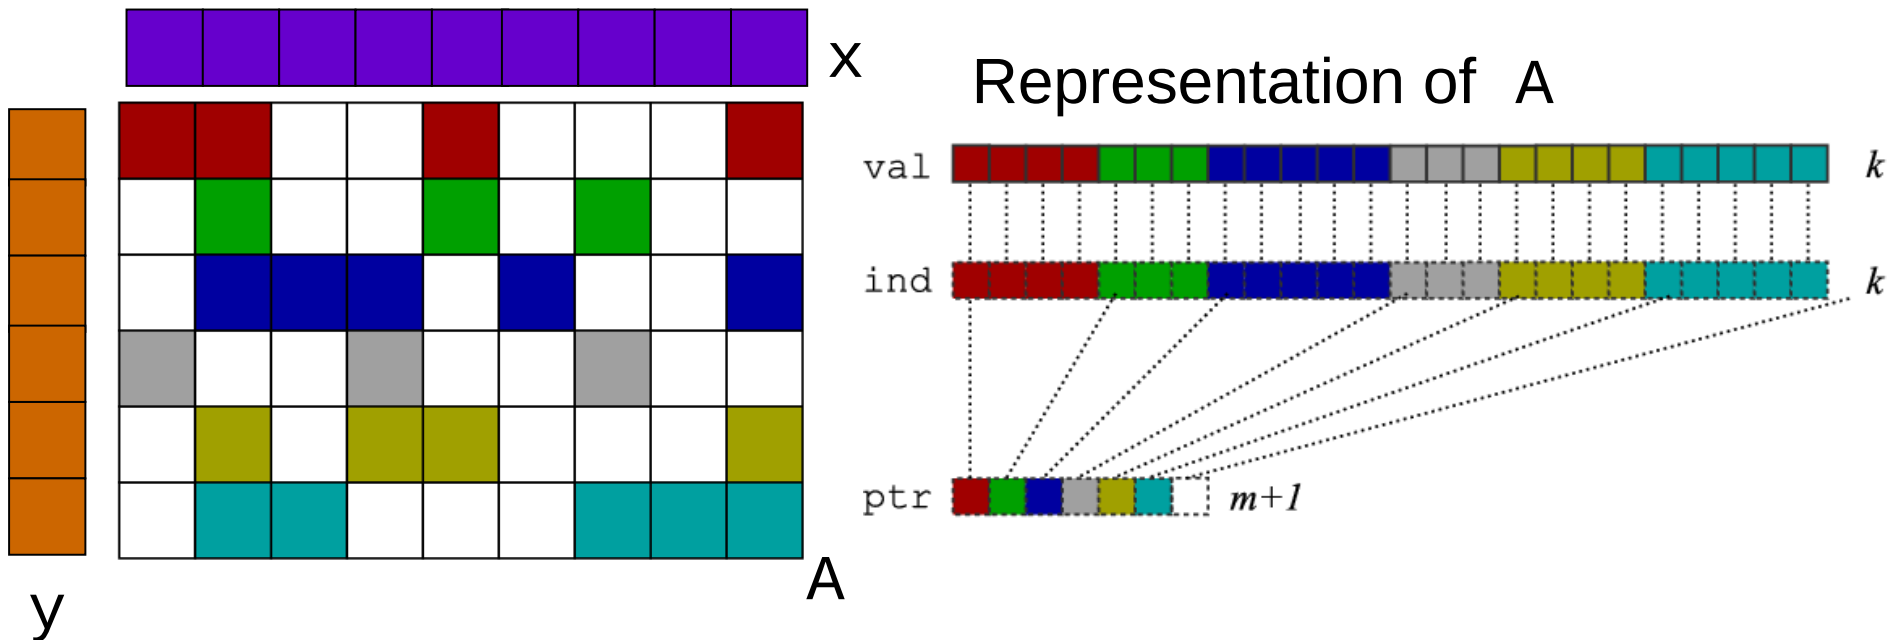
Summary of particle methods

- Model contains discrete entities, namely, particles
- Time is continuous – must be discretized to solve
- Simulation follows particles through timesteps
 - $\text{Force} = \text{external_force} + \text{nearby_force} + \text{far_field_force}$
 - All-pairs algorithm is simple, but inefficient, $O(n^2)$
 - *Particle-mesh* methods approximate by moving particles to a regular mesh, where it is easier to compute forces
 - *Tree-based* algorithms approximate by treating set of particles as a group, when far away

Compressed Sparse Row (CSR) Format

$y = y + A \cdot x$, only store, do arithmetic, on nonzero entries

CSR format is simplest one of many possible data structures for A



Matrix-vector multiply kernel: $y(i) \leftarrow y(i) + A(i,j) \cdot x(j)$

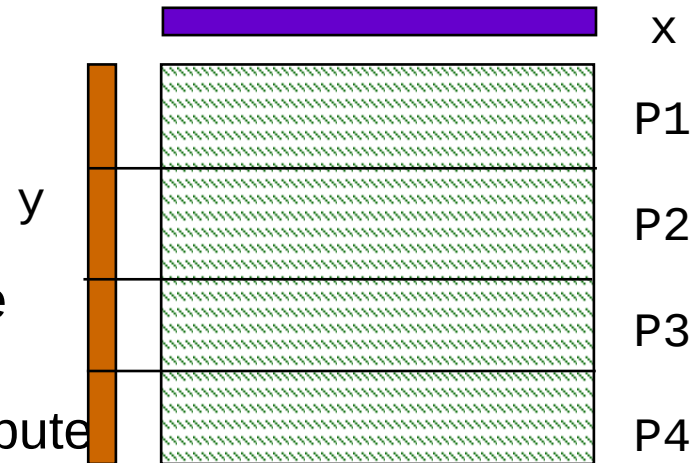
for each row i

for $k=ptr[i]$ to $ptr[i+1]-1$ do

$y[i] = y[i] + val[k] \cdot x[ind[k]]$

Parallel Sparse Matrix-vector multiplication

- $y = A * x$, where A is a sparse $n \times n$ matrix



- Questions
 - which processors store
 - $y[i]$, $x[i]$, and $A[i,j]$
 - which processors compute
 - $y[i] = \text{sum (from 1 to } n) A[i,j] * x[j]$
 $= (\text{row } i \text{ of } A) * x$... a sparse dot product
- Partitioning
 - Partition index set $\{1, \dots, n\} = N1 \cup N2 \cup \dots \cup Np$.
 - For all i in Nk , Processor k stores $y[i]$, $x[i]$, and row i of A
 - For all i in Nk , Processor k computes $y[i] = (\text{row } i \text{ of } A) * x$
 - “owner computes” rule: Processor k compute the $y[i]$ s it owns.

May require communication

Summary of common problems

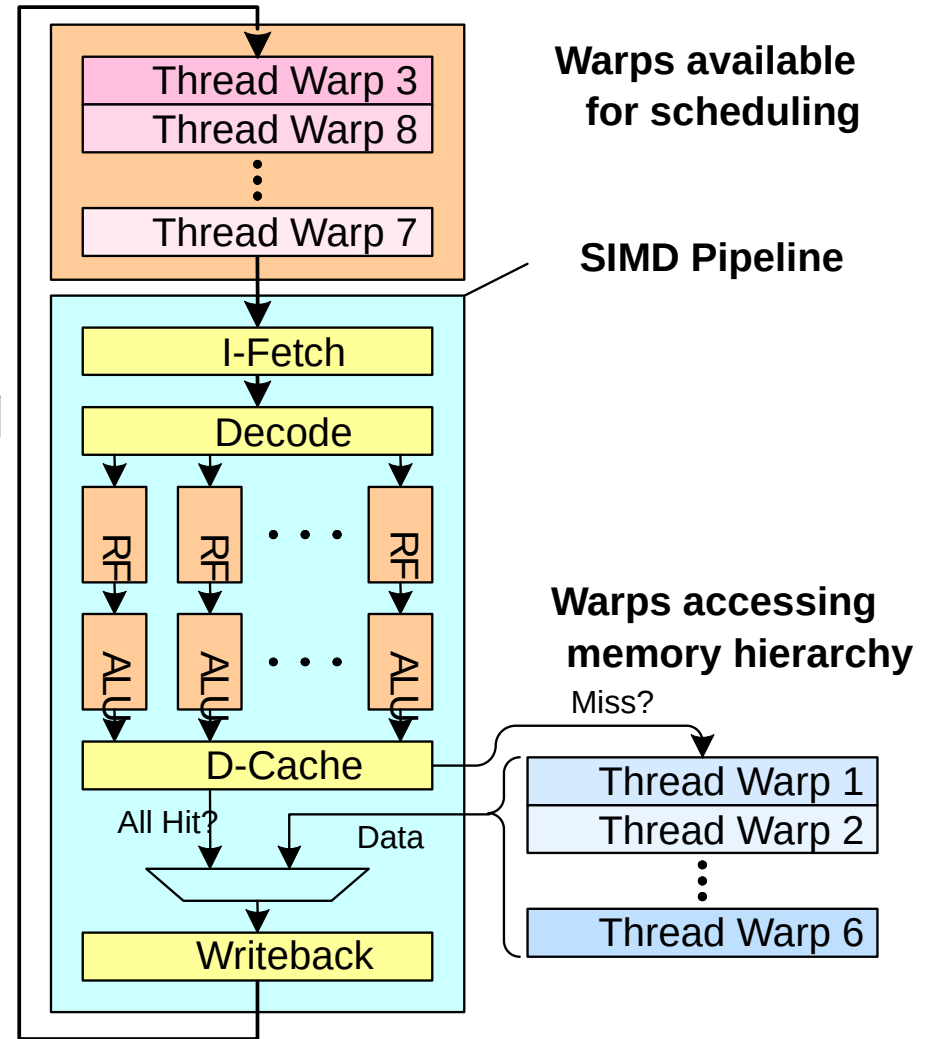
- Load Balancing
 - Dynamically – if load changes significantly during job
 - Statically - Graph partitioning
 - Discrete systems
 - Sparse matrix vector multiplication
- Linear algebra
 - Solving linear systems (sparse and dense)
 - Eigenvalue problems will use similar techniques
- Fast Particle Methods
 - $O(n \log n)$ instead of $O(n^2)$

Summary of tree algorithms

- Lots of problems can be done quickly - in theory - using trees
- Some algorithms are widely used
 - broadcasts, reductions, parallel prefix
 - carry look ahead addition
- Some are of theoretical interest only
 - Csanky's method for matrix inversion
 - Solving tridiagonal linear systems (without pivoting)
 - Both numerically unstable
- Embedded in various systems
 - MPI, NESL (CMU), other languages
 - CM-5 hardware control network

Latency Hiding via Warp-Level Fine Grain Multithreading

- Warp: A set of threads that execute the same instruction (on different data elements)
- Fine-grained multithreading
 - One instruction per thread in pipeline at a time (No interlocking)
 - Interleave warp execution to hide latencies
- Register values of all threads stay in register file
- Fine-Grain multithreading enables long latency tolerance
 - Millions of pixels



Threads

- Threads on desktop CPUs
 - Implemented via lightweight processes (for example)
 - General system scheduler
 - Thrashing when more active threads than processors
- An alternative approach
 - Hardware support for many threads / CPU
 - Modest example: hyperthreading
 - More extreme: Cray MTA-2 and XMT
 - Hide memory latency by thread switching
 - Want many more independent threads than cores
- GPU programming
 - Thread creation / context switching are basically free
 - Want lots of threads (thousands for efficiency?!)

CUDA

- CUDA is a programming model designed for:
 - Heterogeneous architectures
 - Wide SIMD parallelism
 - Scalability
- CUDA provides:
 - A thread abstraction to deal with SIMD
 - Synchronization & data sharing between small thread groups
- CUDA programs are written in C++ with minimal extensions
- OpenCL is inspired by CUDA, but HW & SW vendor neutral

CUDA Threads

- Independent thread of execution
 - has its own program counter, variables (registers), processor state, etc.
 - no implication about how threads are scheduled
- CUDA threads might be **physical** threads
 - as mapped onto GPUs
- CUDA threads might be **virtual** threads
 - might pick 1 block = 1 physical thread on multicore CPU

GPU Conclusions

- GPUs gain efficiency from simpler cores and more parallelism
 - Very wide SIMD (SIMT) for parallel arithmetic and latency-hiding
- Heterogeneous programming with manual offload
 - CPU to run OS, etc. GPU for compute
- Massive (mostly data) parallelism required
 - Memory coalescing helps
- Threads in block share faster memory and barriers
 - Blocks in kernel share slow device memory and atomics

Graph traversal: Depth-first search (DFS)

```
procedure DFS(vertex v)
  v.visited = true
  previsit (v)
  for all v s.t.  $(v, w) \in E$ 
    if (!w.visited) DFS(w)
  postvisit (v)
```

Graph traversal:

Serial Breadth-first search (BFS)

Push seed node onto queue and mark

While queue nonempty:

- Pop node from queue

- Visit node

- Push unmarked neighbors on queue

- Mark all neighbors

Load Balancing Approaches

- Static Load balancing: Divide into equal size parts
 - Straightforward if you know the problem size
 - Can be expensive to estimate / optimize (graph partitioning)
- Dynamic load balancing:
 - When work costs are unknown
 - Observation: Load balancing and locality often trade off

Self Scheduling

- “Self Scheduling” [Tang and Yew ‘86]
 - Shared queue of tasks
 - Processors (workers) take / add tasks to queue

Chunked Scheduling

- “Chunked” (Self) Scheduling [Kruskal and Weiss ‘85]
 - Shared queue of tasks, workers remove chunks of size K
 - Reduce queue contention (locking)
 - How to choose K ?
 - Large chunks: lower contention / overhead
 - Small chunks: even finish time

Guided Self Scheduling

- “Guided” Self Scheduling [Kuck and Polychronopolous ‘87]
 - The chunk size K_i at the i th access to the task pool is given by $K_i = \text{ceiling}(R_i/p)$ where R_i is the # of tasks remaining and p is the number of processors

Not covered

- FFT
- Cloud Computing
- Linear Algebra Applications